

CS162

Operating Systems and Systems Programming

Lecture 7

Concurrency (Continued), Synchronization

February 8th, 2017

Nathan Pemberton

<http://cs162.eecs.Berkeley.edu>

RoadMap


So Far...

- Processes
 - Thread(s) + address space
- Address Space
- Protection
- Dual Mode
- Interrupt handlers
 - Interrupts, exceptions, syscall
- Key Layers: OS Lib, Syscall, Subsystem, Driver
 - User handler on OS descriptors
- Process control
 - fork, wait, signal, exec
- Communication through sockets
 - Integrates processes, protection, file ops, concurrency
- Client-Server Protocol
- Concurrent Execution: Threads

To Come...

- Concurrency/Synchronization (today)
- Filesystems
- Scheduling
- Address Translation/Caching
- Transactions and Distributed Computing
- Security
- ...


Operating System as Design



Perspective on ‘groking’ 162

- Historically, OS was the most complex software
 - Concurrency, synchronization, processes, devices, communication, ...
 - Core systems concepts developed there
- Today, many “applications” are complex software systems too
 - These concepts appear there
 - But they are realized out of the capabilities provided by the operating system
- Seek to understand how these capabilities are implemented upon the basic hardware
- See concepts multiple times from multiple perspectives
 - Lecture provides conceptual framework, integration, examples, ...
 - Book provides a reference with some additional detail
 - Lots of other resources that you need to learn to use
 - » man pages, google, reference manuals, includes (.h)
- Section, Homework and Project provides detail down to the actual code AND direct hands-on experience

Conceptual Framework




- Physical Addresses Shared
 - **So:** Processes and Address Translation
- Single CPU Must Be Shared
 - **So:** Threads
- Processes Aren't Trusted
 - **So:** Kernel/Userspace Split
- Threads Might Not Cooperate
 - **So:** Use timer interrupts to context switch ("preemption")

Recall: MT Kernel IT Process ala Pintos/x86

Kernel


User



- Each user process/thread associated with a kernel thread, described by a 4KB page object containing TCB and kernel stack for the kernel thread

In User thread, w/ Kernel thread waiting

Kernel
User




- x86 CPU holds interrupt SP in register
- During user thread execution, associated kernel thread is “standing by”

User → Kernel

Kernel


User



- Mechanism to resume k-thread goes through interrupt vector


User → Kernel via interrupt vector

Kernel




0

User



255

intr vector




IP
SP
K SP

Proc Regs

PL: 3

- Interrupt transfers control through the Interrupt Vector (IDT in x86)
- iret restores user stack and PL

Pintos Interrupt Processing




Recall: cs61C THE STACK FRAME

Basic Structure of a Function

Prologue

```
entry_label:  
addi $sp,$sp, -framesize  
sw $ra, framesize-4($sp) # save $ra  
save other regs if need be
```



Body... (call other functions...)


Epilogue

```
restore other regs if need be  
lw $ra, framesize-4($sp) # restore $ra  
addi $sp,$sp, framesize  
jr $ra
```




The Stack (review)

- Stack frame includes:
 - Return “instruction” address
 - Parameters
 - Space for other local variables $0xFFFFFFF\rightarrow$
- Stack frames contiguous blocks of memory; stack pointer tells where bottom of stack frame is
- When procedure ends, stack frame is tossed off the stack; frees memory for future stack frames




Pintos Interrupt Processing



In Kernel thread

Kernel

User



- Kernel threads execute with small stack in thread structure
- Scheduler selects among ready kernel and user threads


Timer may trigger thread switch

- `thread_tick`
 - Updates thread counters
 - If quanta exhausted, sets yield flag
- `thread_yield`
 - On path to rtn from interrupt
 - Sets current thread back to READY
 - Pushes it back on ready_list
 - Calls schedule to select next thread to run upon iret
- Schedule
 - Selects next thread to run
 - Calls `switch_threads` to change regs to point to stack for thread to resume
 - Sets its status to RUNNING
 - If user thread, activates the process
 - Returns back to `intr_handler`

Thread Switch (switch.S)

Kernel

User




- `switch_threads`: save regs on current small stack, change SP, return from destination threads call to `switch_threads`


Switch to Kernel Thread for Process

Kernel

User




Pintos Return from Processing



Kernel → User

Kernel

User



- Interrupt return (iret) restores user stack and PL




- Re-visit Slides and Lecture
- Read Book!
- Read Code!
 - e.g. thread.c and thread.S

Goals for Today

- The Concurrency Problem
- Synchronization Operations
- Higher-level Synchronization Abstractions
 - Semaphores, monitors, and condition variables




Recall: Thread Abstraction



- Infinite number of processors
- Threads execute with variable speed
 - Programs must be designed to work with any schedule

Multiprocessing vs Multiprogramming

- Remember Definitions:
 - Multiprocessing ≡ Multiple CPUs or cores or hyperthreads (HW per-instruction interleaving)
 - Multiprogramming ≡ Multiple Jobs or Processes
 - Multithreading ≡ Multiple threads per Process
- What does it mean to run two threads “concurrently”?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...



Correctness for systems with concurrent threads

- If dispatcher can schedule threads in any way, programs must work under all circumstances
 - Can you test for this?
 - How can you know if your program works?
- Independent Threads:
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- Cooperating Threads:
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
 - Sometimes called “Heisenbugs”


Interactions Complicate Debugging

- Is any program truly independent?
 - Every process shares the file system, OS resources, network, etc.
 - Extreme example: buggy device driver causes thread A to crash “independent thread” B
- Non-deterministic errors are really difficult to find
 - Example: Memory layout of kernel+user programs
 - » depends on scheduling, which depends on timer/other things
 - » Original UNIX had a bunch of non-deterministic errors
 - Example: Something which does interesting I/O
 - » User typing of letters used to help generate secure keys

Why allow cooperating threads?


- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - » What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - » Many different file systems do read-ahead
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - More important than you might think
 - Chop large problem up into simpler pieces
 - » To compile, for instance, **gcc** calls **cpp** | **cc1** | **cc2** | **as** | **ld**
 - » Makes system easier to extend

Recall: How does Thread get started?



- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:

```
serverLoop() {  
    connection = AcceptCon();  
    ProcessFork(ServiceWebPage(), connection);  
}
```
- What are some disadvantages of this technique?

Threaded Web Server

- Instead, use a single process
- Multithreaded (cooperating) version:


```
serverLoop() {  
    connection = AcceptCon();  
    ThreadFork(ServiceWebPage(), connection);  
}
```

- Looks almost the same, but has many advantages:
 - Can share file caches kept in memory, results of CGI scripts, other things
 - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- What about Denial of Service attacks or digg / Slashdot effects?



Thread Pools


- Problem with previous version: Unbounded Threads
 - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```
master() {  
    allocThreads(worker, queue);  
    while(TRUE) {  
        con=AcceptCon();  
        Enqueue(queue, con);  
        wakeUp(queue);  
    }  
}
```

```
worker(queue) {  
    while(TRUE) {  
        con=Dequeue(queue);  
        if (con==null)  
            sleepOn(queue);  
        else  
            ServiceWebPage(con);  
    }  
}
```

ATM Bank Server



- ATM server requirements:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}
ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}
Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style
- Example

```
BankServer() {  
    while(TRUE) {  
        event = WaitForNextEvent();  
        if (event == ATMRequest)  
            StartOnRequest();  
        else if (event == AcctAvail)  
            ContinueRequest();  
        else if (event == AcctStored)  
            FinishRequest();  
    }  
}
```

- What if we missed a blocking I/O step?
- What if we have to split code into hundreds of pieces which could be blocking?
- This technique is used for programming GPUs (Graphics Processing Unit)

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without having to “deconstruct” code into non-blocking fragments
 - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);           /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

Thread 1

```
load r1, acct->balance  
  
add r1, amount1  
store r1, acct->balance
```

Thread 2

```
load r1, acct->balance  
add r1, amount2  
store r1, acct->balance
```

Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, what about (Initially, $y = 12$):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?
- Or, what are the possible values of x below?

Thread A

$x = 1;$

Thread B

$x = 2;$

- X could be 1 or 2 (non-deterministic!)
- Could even be 3 for serial processors:
 - » Thread A writes 0001, B writes 0010 → scheduling order ABABABBA yields 3!

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- **Atomic Operation**: an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide can’t happen
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
 - Cooperating threads inherently non-deterministic and non-reproducible
 - Really hard to debug unless carefully designed!
- Examples:




Figure 1. Typical Therac-25 facility

Another Concurrent Program Example

- Two threads, A and B, compete with each other
 - One tries to increment a shared counter
 - The other tries to decrement the counter

<u>Thread A</u>	<u>Thread B</u>
<pre>i = 0; while (i < 10) i = i + 1; printf("A wins!");</pre>	<pre>i = 0; while (i > -10) i = i - 1; printf("B wins!");</pre>

- Assume that memory loads and stores are atomic, but incrementing and decrementing are *not* atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed?
Is it guaranteed that it goes on forever?

Hand Simulation Multiprocessor Example

- Inner loop looks like this:

	<u>Thread A</u>		<u>Thread B</u>
$r1=0$	load $r1, M[i]$		$r1=0$
$r1=1$	add $r1, r1, 1$		$r1=-1$
$M[i]=1$	store $r1, M[i]$		$M[i]=-1$
			store $r1, M[i]$

- Hand Simulation:

- And we're off. A gets off to an early start
- B says “hmph, better go fast” and tries really hard
- A goes ahead and writes “I”
- B goes and writes “-I”
- A says “HUh??? I could have sworn I put a I there”

- Could this happen on a uniprocessor?

- Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...

Administrivia

- Group/Section assignments finalized!
 - If you are not in group, talk to us immediately!
- Attend assigned sections
 - Need to know your TA!
 - » Participation is 5% of your grade
 - » Should attend section with your TA
- First design doc due next **Wednesday**
 - This means you should be well on your way with Project 1
 - Watch for notification from your TA to sign up for design review
- Basic semaphores work in PintOS!
 - However, you will need to implement priority scheduling behavior both in semaphore and ready queue

-
-
- BREAK

Motivation: “Too Much Milk”

- Great thing about OS’s – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:




Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Definitions

- **Synchronization:** using atomic operations to ensure cooperation between threads
 - For now, only loads and stores are atomic
 - We are going to show that it's hard to build anything useful with only reads and writes
- **Mutual Exclusion:** ensuring that only one thread does a particular thing at a time
 - One thread excludes the other while doing its task
- **Critical Section:** piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing

More Definitions

- **Lock:** prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting



- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: roommate angry if only wants OJ



- Of Course – We don't know how to make a lock yet


Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
 - Always write down behavior first
- What are the correctness properties for the “Too much milk” problem???
 - Never more than one person buys
 - Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?
 - Still too much milk **but only occasionally!**
 - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
    }  
}  
remove note;
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk



Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Thread A


```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNoteA) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- Does this work?
- Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
 - *Extremely unlikely* that this would happen, but will at worse possible time
 - Probably something like this in UNIX

Too Much Milk Solution #2: problem!



- I thought you had the milk! But I thought you had the milk!
- This kind of lockup is called “starvation!”

Too Much Milk Solution #3

- Here is a possible two-note solution:

Thread A

```
leave note A;  
while (note B) { //X  
    do nothing;  
}  
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) { //Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit
- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit

Solution #3 discussion

- Our solution protects a single “CriticalSection” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
- There's a better way
 - Have hardware provide better (higher-level) primitives than atomic load and store
 - Build even higher-level programming abstractions on this new hardware support

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock (more in a moment)
 - `Lock.Acquire()` – wait until lock is free, then grab
 - `Lock.Release()` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
milklock.Acquire();  
if (nomilk)  
    buy milk;  
milklock.Release();
```

- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
 - Skip the test since you always need more ice cream.

Where are we going with synchronization?

Programs	Shared Programs			
Higher-level API	Locks	Semaphores	Monitors	Send/Receive
Hardware	Load/Store	Disable Ints	Test&Set	Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

Summary (1 of 2)

- Concurrent threads are a very useful abstraction
 - Allow transparent overlapping of computation and I/O
 - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent
- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives

Next Time

- How to build a lock
- Other synchronization primitives and techniques
- Parallel Programming Paradigms