

CS162 Operating Systems and Systems Programming Lecture 16

General I/O

March 20th, 2017
Prof. Ion Stoica
<http://cs162.eecs.Berkeley.edu>

The Requirements of I/O

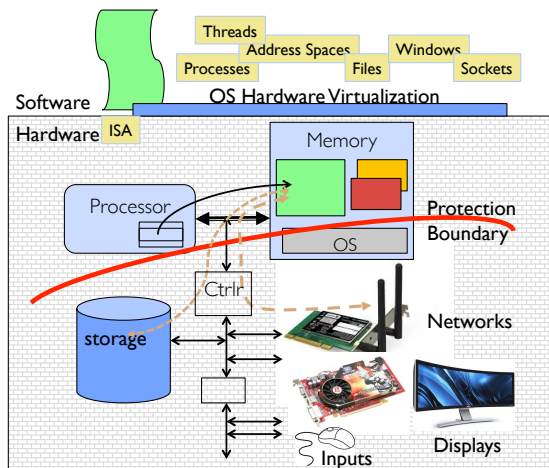
- So far in this course:
 - We have learned how to manage CPU and memory
- What about I/O?
 - Without I/O, computers are useless (disembodied brains?)
 - But... thousands of devices, each slightly different
 - » How can we standardize the interfaces to these devices?
 - Devices unreliable: media failures and transmission errors
 - » How can we make them reliable???
 - Devices unpredictable and/or slow
 - » How can we manage them if we don't know what they will do or how they will perform?

3/20/17

CS162 © UCB Spring 2017

Lec 16.2

OS Basics: I/O

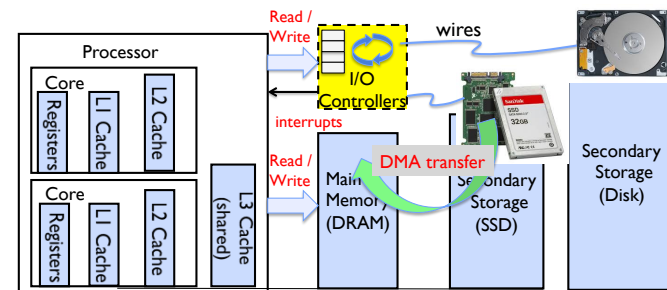


3/20/17

CS162 © UCB Spring 2017

Lec 16.3

In a Picture



- I/O devices you recognize are supported by I/O Controllers
- Processors access them by reading and writing IO registers as if they were memory
 - Write commands and arguments, read status and results

3/20/17

CS162 © UCB Spring 2017

Lec 16.4

Operational Parameters for I/O

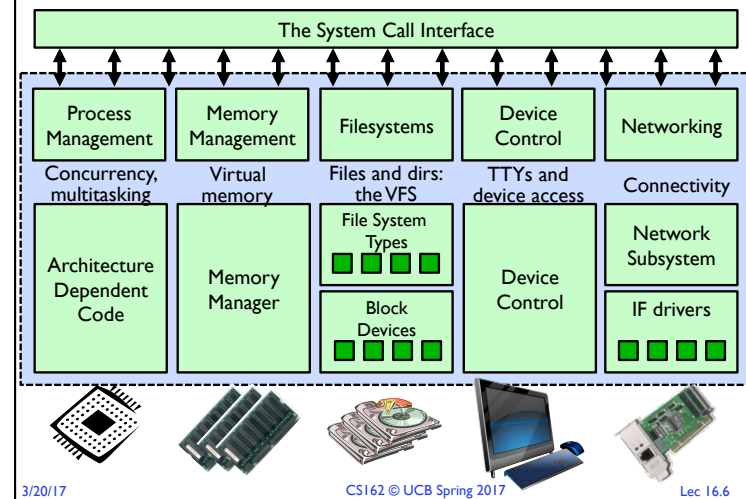
- Data granularity: Byte vs. Block
 - Some devices provide single byte at a time (e.g., keyboard)
 - Others provide whole blocks (e.g., disks, networks, etc.)
- Access pattern: Sequential vs. Random
 - Some devices must be accessed sequentially (e.g., tape)
 - Others can be accessed “randomly” (e.g., disk, cd, etc.)
 - » Fixed overhead to start transfers
 - Some devices require continual monitoring
 - Others generate interrupts when they need service
- Transfer Mechanism: Programmed IO and DMA

3/20/17

CS162 © UCB Spring 2017

Lec 16.5

Kernel Device Structure



3/20/17

CS162 © UCB Spring 2017

Lec 16.6

The Goal of the I/O Subsystem

- Provide Uniform Interfaces, Despite Wide Range of Different Devices
 - This code works on many different devices:


```
FILE fd = fopen("/dev/something", "rw");
for (int i = 0; i < 10; i++) {
    fprintf(fd, "Count %d\n", i);
}
close(fd);
```
 - Why? Because code that controls devices (“device driver”) implements standard interface
- We will try to get a flavor for what is involved in actually controlling devices in rest of lecture
 - Can only scratch surface!

3/20/17

CS162 © UCB Spring 2017

Lec 16.7

Want Standard Interfaces to Devices

- **Block Devices:** e.g. disk drives, tape drives, DVD-ROM
 - Access blocks of data
 - Commands include `open()`, `read()`, `write()`, `seek()`
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character Devices:** e.g. keyboards, mice, serial ports, some USB devices
 - Single characters at a time
 - Commands include `get()`, `put()`
 - Libraries layered on top allow line editing
- **Network Devices:** e.g. Ethernet, Wireless, Bluetooth
 - Different enough from block/character to have own interface
 - Unix and Windows include `socket` interface
 - » Separates network protocol from network operation
 - » Includes `select()` functionality
 - Usage: pipes, FIFOs, streams, queues, mailboxes

3/20/17

CS162 © UCB Spring 2017

Lec 16.8

How Does User Deal with Timing?

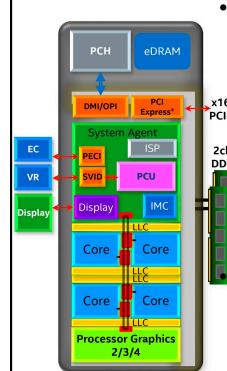
- **Blocking Interface:** "Wait"
 - When request data (e.g. `read()` system call), put process to sleep until data is ready
 - When write data (e.g. `write()` system call), put process to sleep until device is ready for data
- **Non-blocking Interface:** "Don't Wait"
 - Returns quickly from read or write request with count of bytes successfully transferred
 - Read may return nothing, write may write nothing
- **Asynchronous Interface:** "Tell Me Later"
 - When request data, take pointer to user's buffer, return immediately; later kernel fills buffer and notifies user
 - When send data, take pointer to user's buffer, return immediately; later kernel takes data and notifies user

3/20/17

CS162 © UCB Spring 2017

Lec 16.9

Chip-scale Features of 2015 x86 (Sky Lake)

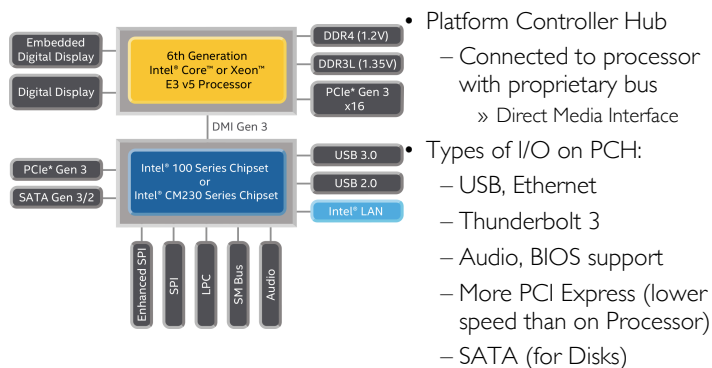
- 
- Significant pieces:
 - Four OOO cores with deeper buffers
 - » New Intel MPX (Memory Protection Extensions)
 - » New Intel SGX (Software Guard Extensions)
 - » Issue up to 6 μ -ops/cycle
 - Integrated GPU, System Agent (Mem, Fast I/O)
 - » Large shared L3 cache with on-chip ring bus
 - » 2 MB/core instead of 1.5 MB/core
 - » High-BW access to L3 Cache
 - Integrated I/O
 - Integrated memory controller (IMC)
 - » Two independent channels of DDR3L/DDR4 DRAM
 - High-speed PCI-Express (for Graphics cards)
 - Direct Media Interface (DMI) Connection to PCH (Platform Control Hub)

3/20/17

CS162 © UCB Spring 2017

Lec 16.10

Sky Lake I/O: PCH



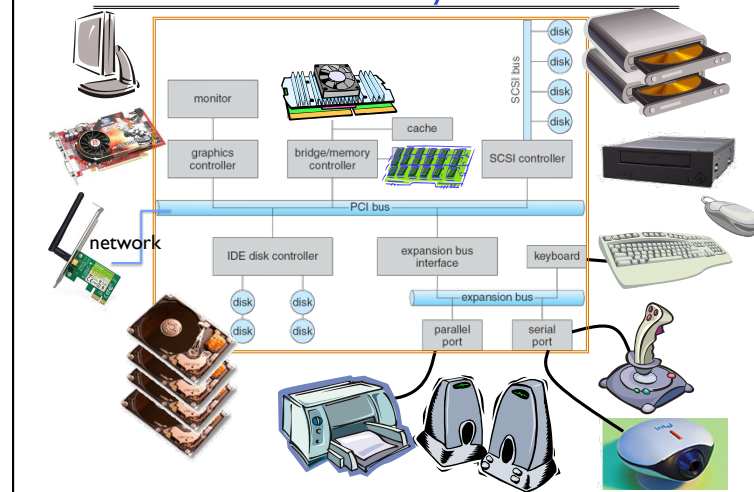
Sky Lake System Configuration

3/20/17

CS162 © UCB Spring 2017

Lec 16.11

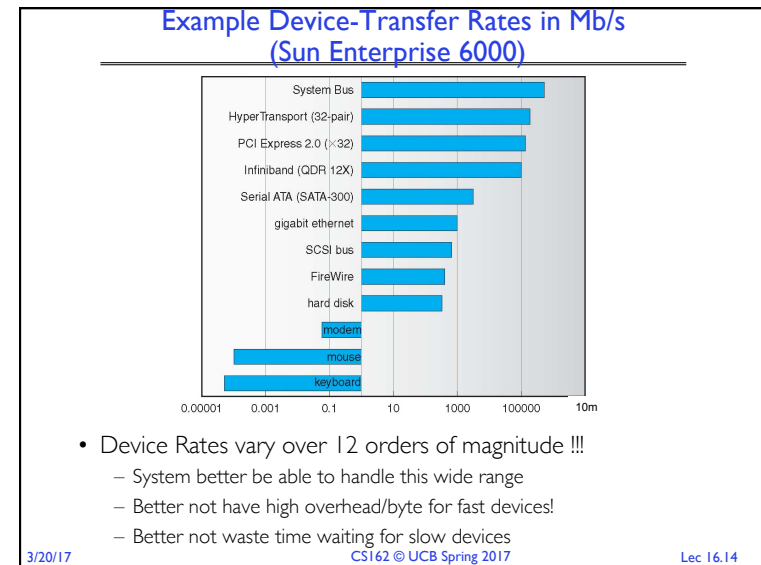
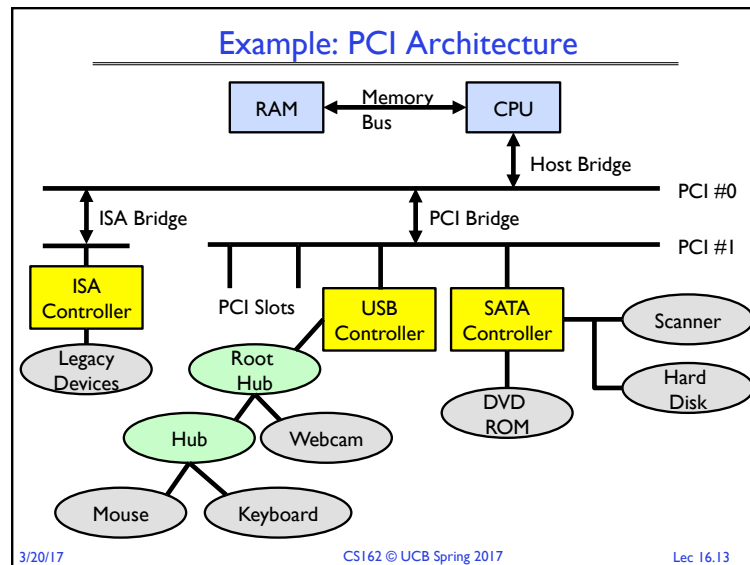
Modern I/O Systems



3/20/17

CS162 © UCB Spring 2017

Lec 16.12



Administrivia

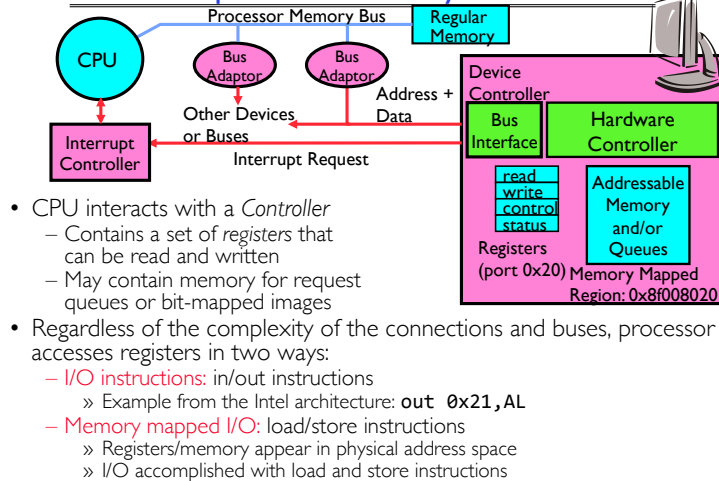
- Midterm 2 **TOMORROW** on **Tue 3/21 7:00-8:30PM**
 - All topics up to and including Lecture 15
 - » Focus will be on Lectures 11 – 15 and associated readings
 - » Projects 1 & 2, Homework 0 – 2
 - Closed book with 2 pages of hand-written notes both sides
 - Room assignments by **last name**:
 - » A-H | 100 Genetics and Plant Biology Building, I-Z | Pimentel

3/20/17 CS162 © UCB Spring 2017 Lec 16.15

BREAK

3/20/17 CS162 © UCB Spring 2017 Lec 16.16

How does the processor actually talk to the device?



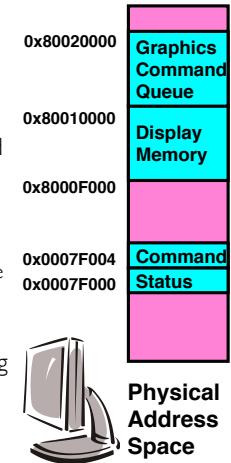
3/20/17

CS162 © UCB Spring 2017

Lec 16.17

Example: Memory-Mapped Display Controller

- Memory-Mapped:
 - Hardware maps control registers and display memory into physical address space
 - Addresses set by HW jumpers or at boot time
 - Simply writing to display memory (also called the "frame buffer") changes image on screen
 - Addr: `0x8000F000` – `0x8000FFFF`
 - Writing graphics description to cmd queue
 - Say enter a set of triangles describing some scene
 - Addr: `0x80010000` – `0x8001FFFF`
 - Writing to the command register may cause on-board graphics hardware to do something
 - Say render the above scene
 - Addr: `0x0007F004`
- Can protect with address translation



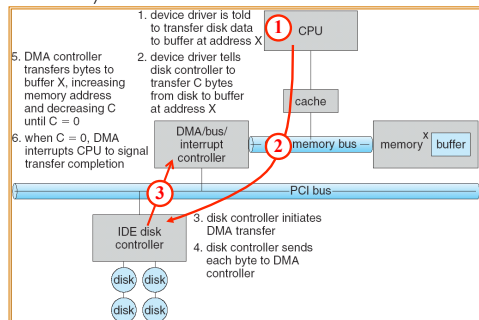
3/20/17

CS162 © UCB Spring 2017

Lec 16.18

Transferring Data To/From Controller

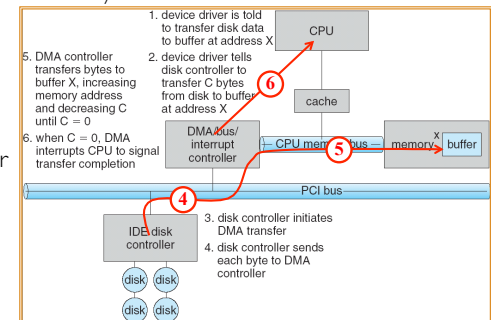
- Programmed I/O:**
 - Each byte transferred via processor in/out or load/store
 - Pro: Simple hardware, easy to program
 - Con: Consumes processor cycles proportional to data size
- Direct Memory Access:**
 - Give controller access to memory bus
 - Ask it to transfer data blocks to/from memory directly
- Sample interaction with DMA controller (from OSC book):



3/20/17

Transferring Data To/From Controller

- Programmed I/O:**
 - Each byte transferred via processor in/out or load/store
 - Pro: Simple hardware, easy to program
 - Con: Consumes processor cycles proportional to data size
- Direct Memory Access:**
 - Give controller access to memory bus
 - Ask it to transfer data blocks to/from memory directly
- Sample interaction with DMA controller (from OSC book):



3/20/17

I/O Device Notifying the OS

- The OS needs to know when:
 - The I/O device has completed an operation
 - The I/O operation has encountered an error
- I/O Interrupt:**
 - Device generates an interrupt whenever it needs service
 - Pro: handles unpredictable events well
 - Con: interrupts relatively high overhead
- Polling:**
 - OS periodically checks a device-specific status register
 - I/O device puts completion information in status register
 - Pro: low overhead
 - Con: may waste many cycles on polling if infrequent or unpredictable I/O operations
- Actual devices combine both polling and interrupts
 - For instance – High-bandwidth network adapter:
 - Interrupt for first incoming packet
 - Poll for following packets until hardware queues are empty

3/20/17

CS162 © UCB Spring 2017

Lec 16.21

Device Drivers

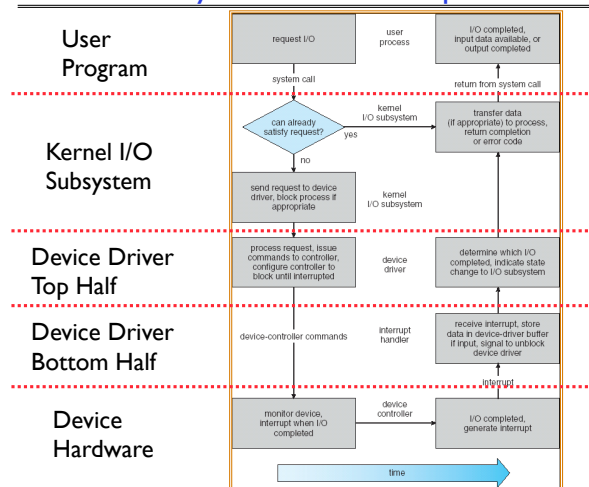
- Device Driver:** Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard, internal interface
 - Same kernel I/O system can interact easily with different device drivers
 - Special device-specific configuration supported with the `ioctl()` system call
- Device Drivers typically divided into two pieces:
 - Top half: accessed in call path from system calls
 - implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
 - This is the kernel's interface to the device driver
 - Top half will start I/O to device, may put thread to sleep until finished
 - Bottom half: run as interrupt routine
 - Gets input or transfers next block of output
 - May wake sleeping threads if I/O now complete

3/20/17

CS162 © UCB Spring 2017

Lec 16.22

Life Cycle of An I/O Request



3/20/17

CS162 © UCB Spring 2017

Lec 16.23

Basic Performance Concepts

- Response Time or Latency:** Time to perform an operation(s)
- Bandwidth or Throughput:** Rate at which operations are performed (op/s)
 - Files: MB/s, Networks: Mb/s, Arithmetic: GFLOP/s
- Start up or "Overhead":** time to initiate an operation
- Most I/O operations are roughly linear in n bytes
 - $\text{Latency}(n) = \text{Overhead} + n/\text{TransferCapacity}$

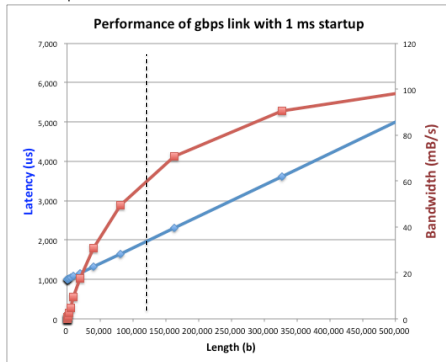
3/20/17

CS162 © UCB Spring 2017

Lec 16.24

Example (Fast Network)

- Consider a 1 Gb/s link ($B = 125 \text{ MB/s}$)
 - With a startup cost $S = 1 \text{ ms}$



- Latency(n) = $S + n/B$
- Bandwidth = $n/(S + n/B) = B*n/(B*S + n) = B/(B*S/n + 1)$

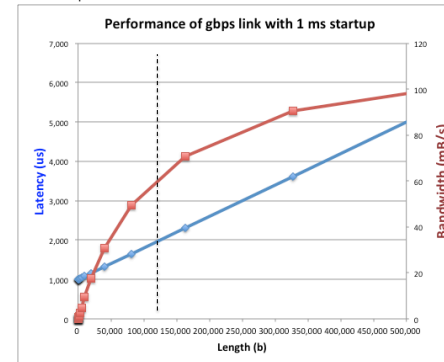
3/20/17

CS162 © UCB Spring 2017

Lec 16.25

Example (Fast Network)

- Consider a 1 Gb/s link ($B = 125 \text{ MB/s}$)
 - With a startup cost $S = 1 \text{ ms}$



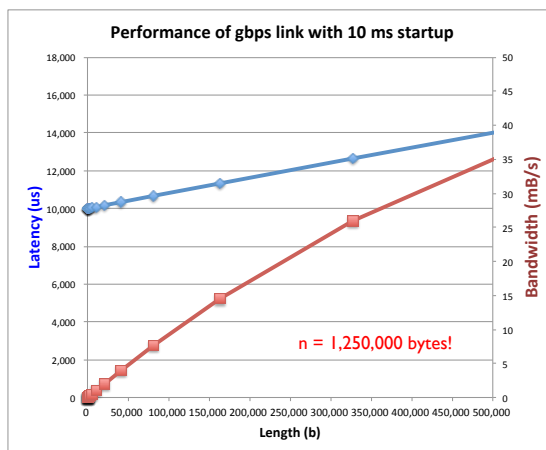
- Bandwidth = $B/(B*S/n + 1)$
- half-power point occurs at $n=S*B \rightarrow \text{Bandwidth} = B/2$

3/20/17

CS162 © UCB Spring 2017

Lec 16.26

Example: at 10 ms startup (like Disk)



3/20/17

CS162 © UCB Spring 2017

Lec 16.27

What Determines Peak BW for I/O ?

- Bus Speed
 - PCI-X: $1064 \text{ MB/s} = 133 \text{ MHz} \times 64 \text{ bit (per lane)}$
 - ULTRA WIDE SCSI: 40 MB/s
 - Serial Attached SCSI & Serial ATA & IEEE 1394 (firewire): 1.6 Gb/s full duplex (200 MB/s)
 - USB 3.0 – 5 Gb/s
 - Thunderbolt 3 – 40 Gb/s
- Device Transfer Bandwidth
 - Rotational speed of disk
 - Write / Read rate of NAND flash
 - Signaling rate of network link
- Whatever is the bottleneck in the path...

3/20/17

CS162 © UCB Spring 2017

Lec 16.28

Storage Devices

- Magnetic disks
 - Storage that rarely becomes corrupted
 - Large capacity at low cost
 - Block level random access (except for SMR – later!)
 - Slow performance for random access
 - Better performance for sequential access
- Flash memory
 - Storage that rarely becomes corrupted
 - Capacity at intermediate cost (5-20x disk)
 - Block level random access
 - Good performance for reads; worse for random writes
 - Erasure requirement in large blocks
 - Wear patterns issue

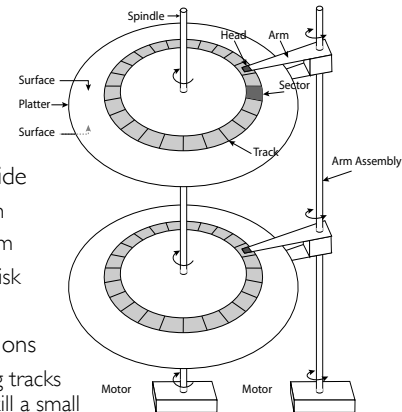
3/20/17

CS162 © UCB Spring 2017

Lec 16.29

The Amazing Magnetic Disk

- Unit of Transfer: Sector
 - Ring of sectors form a track
 - Stack of tracks form a cylinder
 - Heads position on cylinders
- Disk Tracks ~ 1μm (micron) wide
 - Wavelength of light is ~ 0.5μm
 - Resolution of human eye: 50μm
 - 100K tracks on a typical 2.5" disk
- Separated by unused guard regions
 - Reduces likelihood neighboring tracks are corrupted during writes (still a small non-zero chance)



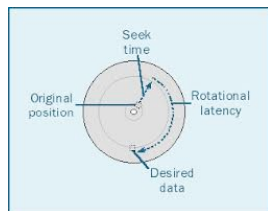
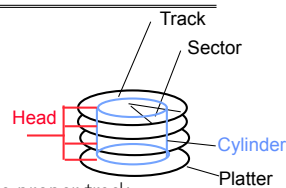
3/20/17

CS162 © UCB Spring 2017

Lec 16.30

Review: Magnetic Disks

- **Cylinders**: all the tracks under the head at a given point on all surface
- Read/write data is a three-stage process:
 - **Seek time**: position the head/arm over the proper track
 - **Rotational latency**: wait for desired sector to rotate under r/w head
 - **Transfer time**: transfer a block of bits (sector) under r/w head



Seek time = 4-8ms
One rotation = 1-2ms
(3600-7200 RPM)

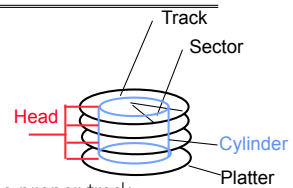
3/20/17

CS162 © UCB Spring 2017

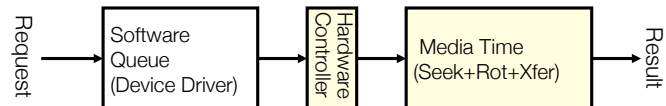
Lec 16.31

Review: Magnetic Disks

- **Cylinders**: all the tracks under the head at a given point on all surface
- Read/write data is a three-stage process:
 - **Seek time**: position the head/arm over the proper track
 - **Rotational latency**: wait for desired sector to rotate under r/w head
 - **Transfer time**: transfer a block of bits (sector) under r/w head



**Disk Latency = Queuing Time + Controller time +
Seek Time + Rotation Time + Xfer Time**



3/20/17

CS162 © UCB Spring 2017

Lec 16.32

Disk Performance Example

- Assumptions:
 - Ignoring queuing and controller times for now
 - Avg seek time of 5ms,
 - 7200RPM \Rightarrow Time for rotation: $60000 \text{ (ms/minute)} / 7200 \text{ (rev/min)} \approx 8 \text{ ms}$
 - Transfer rate of 4MByte/s, sector size of 1 Kbyte \Rightarrow
 $1024 \text{ bytes} / 4 \times 10^6 \text{ (bytes/s)} = 256 \times 10^{-6} \text{ sec} \approx .26 \text{ ms}$
- Read sector from random place on disk:
 - Seek (5ms) + Rot. Delay (4ms) + Transfer (0.26ms)
 - Approx 10ms to fetch/put data: **100 KByte/sec**
- Read sector from random place in same cylinder:
 - Rot. Delay (4ms) + Transfer (0.26ms)
 - Approx 5ms to fetch/put data: **200 KByte/sec**
- Read next sector on same track:
 - Transfer (0.26ms): **4 MByte/sec**
- Key to using disk effectively (especially for file systems) is
to minimize seek and rotational delays

3/20/17

CS162 © UCB Spring 2017

Lec 16.33

Typical Numbers for Magnetic Disk

Parameter	Info / Range
Space/Density	Space: 8TB (Seagate), 10TB (Hitachi) in 3½ inch form factor! Areal Density: ≥ 1 Terabit/square inch! (SMR, Helium, ...)
Average seek time	Typically 5-10 milliseconds. Depending on reference locality, actual cost may be 25-33% of this number.
Average rotational latency	Most laptop/desktop disks rotate at 3600-7200 RPM (16-8 ms/rotation). Server disks up to 15,000 RPM. Average latency is halfway around disk so 8-4 milliseconds
Controller time	Depends on controller hardware
Transfer time	Typically 50 to 100 MB/s. Depends on: <ul style="list-style-type: none"> Transfer size (usually a sector): 512B – 1KB per sector Rotation speed: 3600 RPM to 15000 RPM Recording density: bits per inch on a track Diameter: ranges from 1 in to 5.25 in
Cost	Used to drop by a factor of two every 1.5 years (or even faster); now slowing down

3/20/17

CS162 © UCB Spring 2017

Lec 16.34

(Lots of) Intelligence in the Controller

- Sectors contain sophisticated error correcting codes
 - Disk head magnet has a field wider than track
 - Hide corruptions due to neighboring track writes
- Sector sparing
 - Remap bad sectors transparently to spare sectors on the same surface
- Slip sparing
 - Remap all sectors (when there is a bad sector) to preserve sequential behavior
- Track skewing
 - Sector numbers offset from one track to the next, to allow for disk head movement for sequential ops
- ...

3/20/17

CS162 © UCB Spring 2017

Lec 16.35

Seagate Enterprise

10 TB (2016)

- 7 platters, 14 heads
- 7200 RPMs
- 6 Gbps SATA / 12Gbps SAS interface
- 220MB/s transfer rate, cache size: 256MB
- Helium filled: reduce friction and power usage
- Price: \$500 (\$0.05/GB)



IBM Personal Computer/AT (1986)

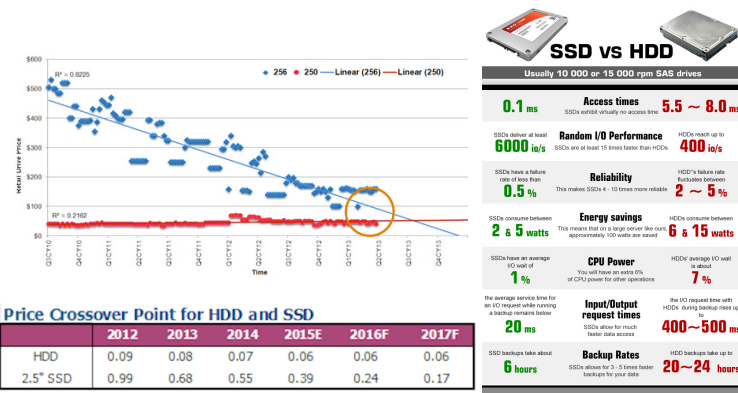
- 30 MB hard disk
- 30-40ms seek time
- 0.7-1 MB/s (est.)
- Price: \$500 (\$17K/GB, 340,000x more expensive !!)

3/20/17

CS162 © UCB Spring 2017

Lec 16.36

HDD vs SSD Comparison



SSD prices drop much faster than HDD

3/20/17

CS162 © UCB Spring 2017

Lec 16.37

Largest SSDs

- 60TB (2016)
- Dual port: 16Gbs
- Seq reads: 1.5GB/s
- Seq writes: 1GB/s
- Random Read Ops (IOPS): 150K
- Price: ~ \$20K (\$0.33/GB)



3/20/17

CS162 © UCB Spring 2017

Lec 16.38

Summary

- I/O Devices Types:
 - Many different speeds (0.1 bytes/sec to GBytes/sec)
 - Different Access Patterns:
 - » Block Devices, Character Devices, Network Devices
 - Different Access Timing:
 - » Blocking, Non-blocking, Asynchronous
- I/O Controllers: Hardware that controls actual device
 - Processor Accesses through I/O instructions, load/store to special physical memory
- Notification mechanisms
 - Interrupts
 - Polling: Report results through status register that processor looks at periodically
- Device drivers interface to I/O devices
 - Provide clean Read/Write interface to OS above
 - Manipulate devices through PIO, DMA & interrupt handling
 - Three types: block, character, and network

3/20/17

CS162 © UCB Spring 2017

Lec 16.39