

# CS162 Operating Systems and Systems Programming Lecture 18

## File Systems

April 3rd, 2017  
Prof. Ion Stoica  
<http://cs162.eecs.Berkeley.edu>

## Recall: How do we Hide I/O Latency?

- **Blocking Interface:** "Wait"
  - When request data (e.g., read() system call), put process to sleep until data is ready
  - When write data (e.g., write() system call), put process to sleep until device is ready for data
- **Non-blocking Interface:** "Don't Wait"
  - Returns quickly from read or write request with count of bytes successfully transferred to kernel
  - Read may return nothing, write may write nothing
- **Asynchronous Interface:** "Tell Me Later"
  - When requesting data, take pointer to user's buffer, return immediately; later kernel fills buffer and notifies user
  - When sending data, take pointer to user's buffer, return immediately; later kernel takes data and notifies user

4/3/17

CS162 ©UCB Spring 2017

Lec 18.2

## I/O & Storage Layers

### Operations, Entities and Interface



4/3/17

CS162 ©UCB Spring 2017

Lec 18.3

## Recall: C Low level I/O

- Operations on File Descriptors – as OS object representing the state of a file
  - User has a "handle" on the descriptor

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int open (const char *filename, int flags [, mode_t mode])
int create (const char *filename, mode_t mode)
int close (int filedes)
```

Bit vector of:

- Access modes (Rd, Wr, ...)
- Open Flags (Create, ...)
- Operating modes (Append, ...)

Bit vector of Permission Bits:

- User|Group|Other X R|W|X

[http://www.gnu.org/software/libc/manual/html\\_node/Opening-and-Closing-Files.html](http://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html)

4/3/17

CS162 ©UCB Spring 2017

Lec 18.4

## Recall: C Low Level Operations

```
ssize_t read (int filedes, void *buffer, size_t maxsize)
- returns bytes read, 0 => EOF, -1 => error
ssize_t write (int filedes, const void *buffer, size_t size)
- returns bytes written
off_t lseek (int filedes, off_t offset, int whence)
- set the file offset
  * if whence == SEEK_SET: set file offset to "offset"
  * if whence == SEEK_CUR: set file offset to crt location + "offset"
  * if whence == SEEK_END: set file offset to file size + "offset"
int fsync (int fildes)
- wait for i/o of filedes to finish and commit to disk
void sync (void) - wait for ALL to finish and commit to disk
```

- When write returns, data is on its way to disk and can be read, but it may not actually be permanent!

4/3/17

CS162 @UCB Spring 2017

Lec 18.5

## Building a File System

- **File System:** Layer of OS that transforms block interface of disks (or other block devices) into Files, Directories, etc.
- File System Components
  - **Naming:** Interface to find files by name, not by blocks
  - **Disk Management:** collecting disk blocks into files
  - **Protection:** Layers to keep data secure
  - **Reliability/Durability:** Keeping of files durable despite crashes, media failures, attacks, etc.

4/3/17

CS162 @UCB Spring 2017

Lec 18.6

## Recall: User vs. System View of a File


- User's view:
  - Durable Data Structures
- System's view (system call interface):
  - Collection of Bytes (UNIX)
  - Doesn't matter to system what kind of data structures you want to store on disk!
- System's view (inside OS):
  - Collection of blocks (a block is a logical transfer unit, while a sector is the physical transfer unit)
  - Block size  $\geq$  sector size; in UNIX, block size is 4KB

4/3/17

CS162 @UCB Spring 2017

Lec 18.7

## Recall: Translating from User to System View



- What happens if user says: give me bytes 2—12?
  - Fetch block corresponding to those bytes
  - Return just the correct portion of the block
- What about: write bytes 2—12?
  - Fetch block
  - Modify portion
  - Write out Block
- Everything inside File System is in whole size blocks
  - For example, **getc()**, **putc()**  $\Rightarrow$  buffers something like 4096 bytes, even if interface is one byte at a time
- From now on, file is a collection of blocks

4/3/17

CS162 @UCB Spring 2017

Lec 18.8

## Disk Management Policies (1/2)

- Basic entities on a disk:
  - **File**: user-visible group of blocks arranged sequentially in logical space
  - **Directory**: user-visible index mapping names to files
- Access disk as linear array of sectors. Two Options:
  - Identify sectors as vectors [cylinder, surface, sector], sort in cylinder-major order, not used anymore
  - **Logical Block Addressing (LBA)**: Every sector has integer address from zero up to max number of sectors
  - Controller translates from address  $\Rightarrow$  physical position
    - » First case: OS/BIOS must deal with bad sectors
    - » Second case: hardware shields OS from structure of disk

4/3/17

CS162 @UCB Spring 2017

Lec 18.9

## Recall: Disk Management Policies (2/2)

- Need way to track free disk blocks
  - Link free blocks together  $\Rightarrow$  too slow today
  - Use bitmap to represent free space on disk
- Need way to structure files: **File Header**
  - Track which blocks belong at which offsets within the logical file structure
  - **Optimize placement of files' disk blocks to match access and usage patterns**

4/3/17

CS162 @UCB Spring 2017

Lec 18.10

## Designing a File System ...


- What factors are critical to the design choices?
- Durable data store  $\Rightarrow$  it's all on disk
- (Hard) Disks Performance !!!
  - Maximize sequential access, minimize seeks
- Open before Read/Write
  - Can perform protection checks and look up where the actual file resource are, in advance
- Size is determined as they are used !!!
  - Can write (or read zeros) to expand the file
  - Start small and grow, need to make room
- Organized into directories
  - What data structure (on disk) for that?
- Need to allocate / free blocks
  - Such that access remains efficient

4/3/17

CS162 @UCB Spring 2017

Lec 18.11

## Components of a File System



4/3/17

CS162 @UCB Spring 2017

Lec 18.12

## Components of a file system

file name  $\xrightarrow{\text{offset}}$  file number  $\xrightarrow{\text{offset}}$  Storage block  
 offset directory offset index structure


- Open performs *Name Resolution*
  - Translates pathname into a “file number”
    - » Used as an “index” to locate the blocks
  - Creates a file descriptor in PCB within kernel
  - Returns a “handle” (another integer) to user process
- Read, Write, Seek, and Sync operate on handle
  - Mapped to file descriptor and to blocks

4/3/17

CS162 ©UCB Spring 2017

Lec 18.13

## Directories



4/3/17

CS162 ©UCB Spring 2017

Lec 18.14

## Directory

- Basically a hierarchical structure
- Each directory entry is a collection of
  - Files
  - Directories
    - » A link to another entries
- Each has a name and attributes
  - Files have data
- Links (hard links) make it a DAG, not just a tree
  - Softlinks (aliases) are another name for an entry


4/3/17

CS162 ©UCB Spring 2017

Lec 18.15

## I/O & Storage Layers

Application / Service



streams

handles


registers

descriptors

Commands and Data Transfers

Disks, Flash, Controllers, DMA

#4 - handle



Directory Structure

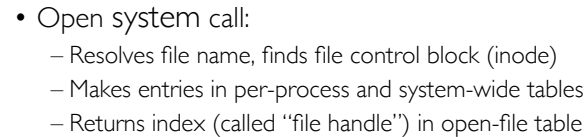
4/3/17

CS162 ©UCB Spring 2017

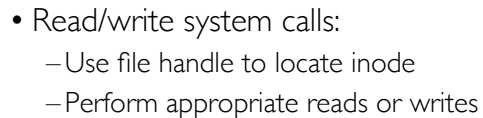
Lec 18.16

- 
- The diagram illustrates the internal structure of a file. On the left, a horizontal blue line represents the **File handle**. Below it, the text **File descriptor** is shown. To the right of the file handle, a vertical stack of blue rectangles represents **Data blocks**. Below the data blocks, a small blue rectangle represents the **Fileobject (inode)**. To the right of the fileobject, a blue cylinder represents the **Position**.

Lec 18.17



Lec 18.18




Lec 18.19

- 
- The diagram shows the mapping of file blocks to disk blocks using a File Allocation Table (FAT). On the left, a vertical column represents the FAT, with entries indexed from 0 to N-1. A blue arrow points from the text 'file number' to the entry at index 31. The entry at index 31 is highlighted in green and contains the value 31. Below it, another green entry is highlighted, and a blue arrow points from the text 'Block 2' to it. On the right, a vertical column represents the Disk Blocks, also indexed from 0 to N-1. The entry at index 31 is highlighted in green and contains the text 'File 31, Block 0'. The entry at index 32 is highlighted in green and contains the text 'File 31, Block 1'. The entry at index 33 is highlighted in green and contains the text 'File 31, Block 2'. Blue arrows indicate the mapping: from the FAT entry at index 31 to the Disk Block entry at index 31, and from the FAT entry at index 32 to the Disk Block entry at index 32. A blue cylinder representing a disk is shown at the bottom right.

Lec 18.20

## FAT Properties

- File is collection of disk blocks
- FAT is linked list I-I with blocks
- File Number is index of root of block list for the file
- File offset ( $o = \langle B, x \rangle$ )
- Follow list to get block #
- Unused blocks  $\Leftrightarrow$  FAT free list



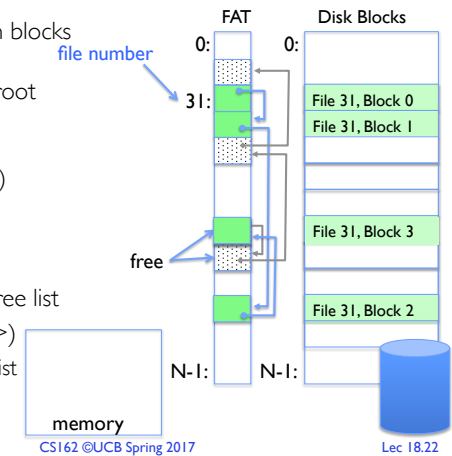
4/3/17

CS162 @UCB Spring 2017

Lec 18.21

## FAT Properties

- File is collection of disk blocks
- FAT is linked list I-I with blocks
- File Number is index of root of block list for the file
- File offset ( $o = \langle B, x \rangle$ )
- Follow list to get block #
- Unused blocks  $\Leftrightarrow$  FAT free list
- Ex: file\_write(31,  $\langle 3, y \rangle$ )
  - Grab blocks from free list
  - Linking them into file




4/3/17

CS162 @UCB Spring 2017

Lec 18.22

## FAT Properties

- File is collection of disk blocks
- FAT is linked list I-I with blocks
- File Number is index of root of block list for the file
- Grow file by allocating free blocks and linking them in
- Ex: Create file, write, write



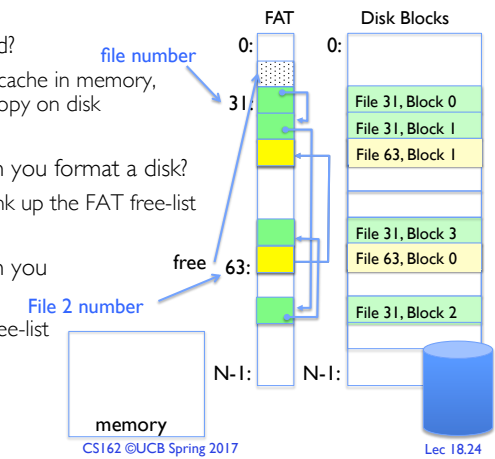
4/3/17

CS162 @UCB Spring 2017

Lec 18.23

## FAT Assessment

- FAT32 (32 instead of 12 bits) used in Windows, USB drives, SD cards, ...**
- Where is FAT stored?
  - On Disk, on boot cache in memory, second (backup) copy on disk
- What happens when you format a disk?
  - Zero the blocks, link up the FAT free-list
- What happens when you quick format a disk?
  - Link up the FAT free-list
- Simple**
  - Can implement in device firmware




4/3/17

CS162 @UCB Spring 2017

Lec 18.24

## FAT Assessment – Issues

- Time to find block (large files) ??
- Block layout for file ???
- Sequential Access ???
- Random Access ???
- Fragmentation ???
  - MSDOS defrag tool
- Small files ???
- Big files ???




4/3/17

CS162 ©UCB Spring 2017

Lec 18.25

## What about the Directory?



- Essentially a file containing  
 <file\_name: file\_number> mappings
- Free space for new entries
- In FAT: file attributes are kept in directory (!!!)
- Each directory a linked list of entries
- Where do you find root directory ( "/" )?

4/3/17

CS162 ©UCB Spring 2017

Lec 18.26

## Directory Structure (cont'd)

- How many disk accesses to resolve "/my/book/count"?
  - Read in file header for root (fixed spot on disk)
  - Read in first data block for root
    - » Table of file name/index pairs. Search linearly – ok since directories typically very small
  - Read in file header for "my"
  - Read in first data block for "my"; search for "book"
  - Read in file header for "book"
  - Read in first data block for "book"; search for "count"
  - Read in file header for "count"
- **Current working directory:** Per-address-space pointer to a directory (inode) used for resolving file names
  - Allows user to specify relative filename instead of absolute path (say CWD="/my/book" can resolve "count")

4/3/17

CS162 ©UCB Spring 2017

Lec 18.27

## Many Huge FAT Security Holes!

- FAT has no access rights
- FAT has no header in the file blocks
- Just gives an index into the FAT
  - (file number = block number)

4/3/17

CS162 ©UCB Spring 2017

Lec 18.28

## Administrivia

- Project 2
  - Code due **today** Monday 4/3
  - Final report due Wednesday 4/5

4/3/17

CS162 @UCB Spring 2017

Lec 18.29

BREAK

4/3/17

CS162 @UCB Spring 2017


Lec 18.30

## Characteristics of Files

### A Five-Year Study of File-System Metadata

NITIN AGRAWAL  
University of Wisconsin, Madison  
and  
WILLIAM J. BOLOSKY, JOHN R. DOUCEUR, and JACOB R. LORCH  
Microsoft Research

9/9




4/3/17

CS162 @UCB Spring 2017

Lec 18.31

## Characteristics of Files

- Most files are small, growing numbers of files over



4/3/17

CS162 @UCB Spring 2017

Lec 18.32



## Characteristics of Files

- Most of the space is occupied by the rare big ones




Fig. 4. Histograms of bytes by containing file size.


4/3/17

CS162 @UCB Spring 2017

Lec 18.33

## So What About a “Real” File System?

- Meet the inode:



4/3/17

CS162 @UCB Spring 2017

Lec 18.34

## An “Almost Real” File System

- Pintos: `src/filesys/file.c`, `inode.c`

```

/* An open file. */
struct file
{
    struct inode *inode; /* File's inode. */
    off_t pos; /* Current position. */
    bool deny_write; /* Has file_deny_write() been called? */
};

/* In-memory inode. */
struct inode
{
    struct list_elem elem; /* Element in inode list. */
    block_sector_t sector; /* Sector number of disk location. */
    int open_cnt; /* Number of openers. */
    bool removed; /* True if deleted, false otherwise. */
    int deny_write_cnt; /* 0: writes ok, >0: deny writes. */
    struct inode_disk data; /* Inode content. */
};

/* On-disk inode.
   Must be exactly BLOCK_SECTOR_SIZE bytes long. */
struct inode_disk
{
    block_sector_t start; /* First data sector. */
    off_t length; /* File size in bytes. */
    unsigned magic; /* Magic number. */
    uint32_t unused[125]; /* Not used. */
};
    
```

4/3/17

## Unix File System

- Original inode format appeared in BSD 4.1
  - Berkeley Standard Distribution Unix
  - Part of your heritage!
  - Similar structure for Linux Ext2/3
- File Number is index into inode arrays
- Multi-level index structure
  - Great for little and large files
  - Asymmetric tree with fixed sized blocks
- Metadata associated with the file
  - Rather than in the directory that points to it
- UNIX Fast File System (FFS) BSD 4.2 Locality Heuristics:
  - Block group placement
  - Reserve space
- Scalable directory structure

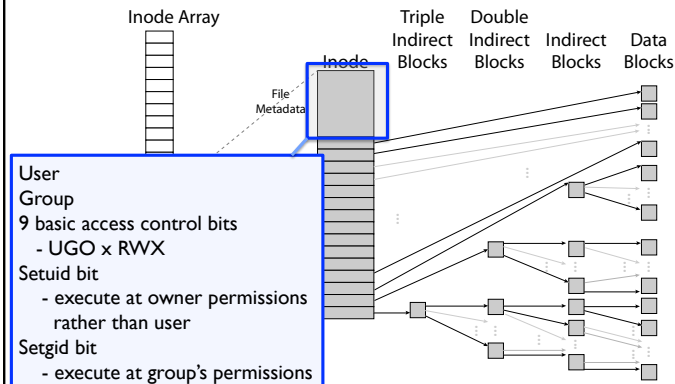
4/3/17

CS162 @UCB Spring 2017

Lec 18.36

## File Attributes

- inode metadata



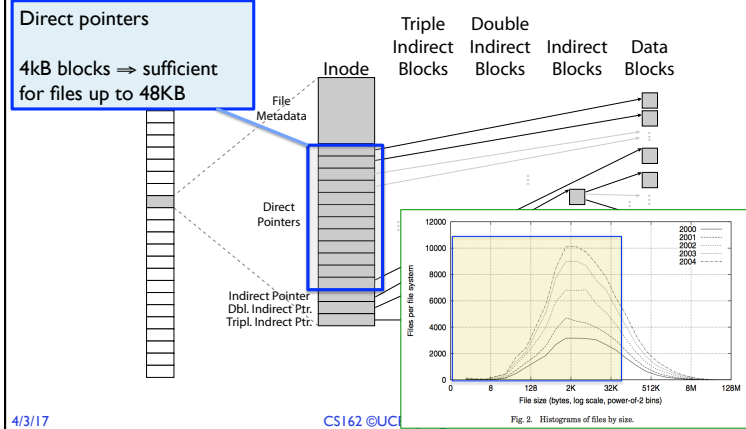
4/3/17

CS162 @UCB Spring 2017

Lec 18.37

## Data Storage

- Small files: 12 pointers direct to data blocks

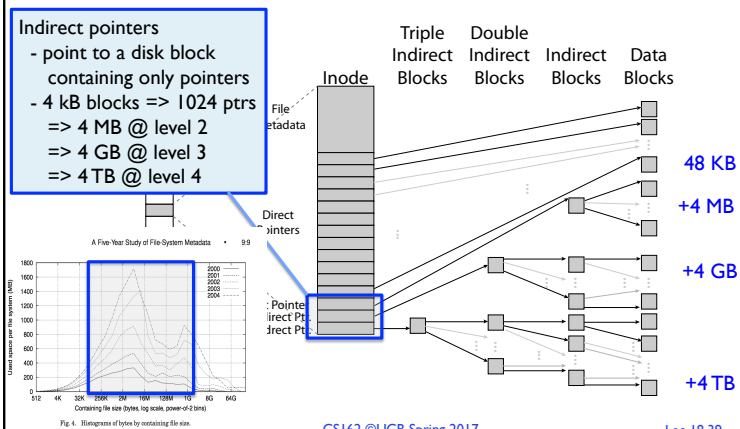


4/3/17

CS162 @UCB

## Data Storage

- Large files: 1,2,3 level indirect pointers



4/3/17

CS162 @UCB Spring 2017

Lec 18.39

## Summary

- File System:
  - Transforms blocks into Files and Directories
  - Optimize for access and usage patterns
  - Maximize sequential access, allow efficient random access
- File (and directory) defined by header, called "inode"
- File Allocation Table (FAT) Scheme
  - Linked-list approach
  - Very widely used: Cameras, USB drives, SD cards
  - Simple to implement, but poor performance and no security
- Look at actual file access patterns – many small files, but large files take up all the space!

4/3/17

CS162 @UCB Spring 2017

Lec 18.40