

CSI 62  
Operating Systems and  
Systems Programming  
Lecture 23

RPC,  
Key-Value Stores,  
Chord

April 19<sup>th</sup>, 2017  
Prof. Ion Stoica  
<http://cs162.eecs.Berkeley.edu>

## Remote Procedure Call (RPC)

- Raw messaging is a bit too low-level for programming
  - Must wrap up information into message at source
  - Must decide what to do with message at destination
  - May need to sit and wait for multiple messages to arrive
- Another option: Remote Procedure Call (RPC)
  - Calls a procedure on a remote machine
  - Client calls:  
`remoteFileSystem→Read("rutabaga");`
  - Translated automatically into call on server:  
`fileSys→Read("rutabaga");`

4/19/17

CS162 ©UCB Spring 2017

Lec 23.2

## RPC Implementation

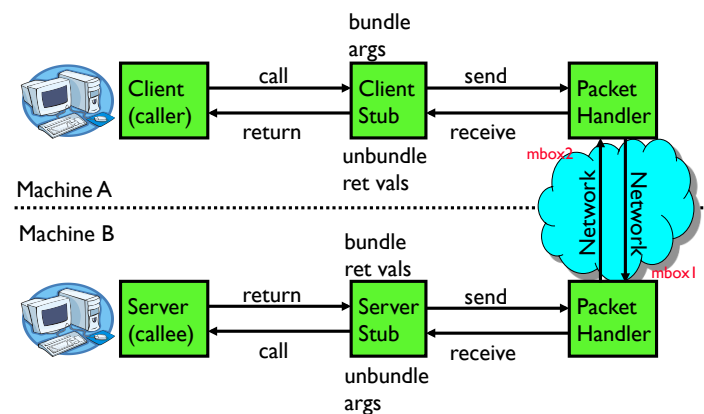
- Request-response message passing (under covers!)
- “Stub” provides glue on client/server
  - Client stub is responsible for “marshalling” arguments and “unmarshalling” the return values
  - Server-side stub is responsible for “unmarshalling” arguments and “marshalling” the return values.
- **Marshalling** involves (depending on system)
  - Converting values to a canonical form, serializing objects, copying arguments passed by reference, etc.

4/19/17

CS162 ©UCB Spring 2017

Lec 23.3

## RPC Information Flow



4/19/17

CS162 ©UCB Spring 2017

Lec 23.4

### RPC Details (1/3)

- Equivalence with regular procedure call
  - Parameters  $\leftrightarrow$  Request Message
  - Result  $\leftrightarrow$  Reply message
  - Name of Procedure: Passed in request message
  - Return Address: mbox2 (client return mail box)
- Stub generator: Compiler that generates stubs
  - Input: interface definitions in an “interface definition language (IDL)”
    - » Contains, among other things, types of arguments/return
  - Output: stub code in the appropriate source language
    - » Code for client to pack message, send it off, wait for result, unpack result and return to caller
    - » Code for server to unpack message, call procedure, pack results, send them off

4/19/17

CS162 ©UCB Spring 2017

Lec 23.5

### RPC Details (2/3)

- Cross-platform issues:
  - What if client/server machines are different architectures/ languages?
    - » Convert everything to/from some canonical form
    - » Tag every item with an indication of how it is encoded (avoids unnecessary conversions)
- How does client know which mbox to send to?
  - Need to translate name of remote service into network endpoint (Remote machine, port, possibly other info)
  - **Binding**: the process of converting a user-visible name into a network endpoint
    - » This is another word for “naming” at network level
    - » Static: fixed at compile time
    - » Dynamic: performed at runtime

4/19/17

CS162 ©UCB Spring 2017

Lec 23.6

### RPC Details (3/3)

- Dynamic Binding
  - Most RPC systems use dynamic binding via name service
    - » Name service provides dynamic translation of service  $\rightarrow$  mbox
  - Why dynamic binding?
    - » Access control: check who is permitted to access service
    - » Fail-over: If server fails, use a different one
- What if there are multiple servers?
  - Could give flexibility at binding time
    - » Choose unloaded server for each new client
  - Could provide same mbox (router level redirect)
    - » Choose unloaded server for each new request
    - » Only works if no state carried from one call to next
- What if multiple clients?
  - Pass pointer to client-specific return mbox in request

4/19/17

CS162 ©UCB Spring 2017

Lec 23.7

### Problems with RPC: Non-Atomic Failures

- Different failure modes in dist. system than on a single machine
- Consider many different types of failures
  - User-level bug causes address space to crash
  - Machine failure, kernel bug causes all processes on same machine to fail
  - Some machine is compromised by malicious party
- Before RPC: whole system would crash/die
- After RPC: One machine crashes/compromised while others keep working
- Can easily result in inconsistent view of the world
  - Did my cached data get written back or not?
  - Did server do what I requested or not?
- Answer? Distributed transactions/Byzantine Commit

4/19/17

CS162 ©UCB Spring 2017

Lec 23.8

### Problems with RPC: Performance

---

- Cost of Procedure call « same-machine RPC « network RPC
- Means programmers must be aware that RPC is not free
  - Caching can help, but may make failure handling complex

4/19/17

CS162 ©UCB Spring 2017

Lec 23.9

### Administrivia

---

- Midterm 3 coming up on **Mon 4/24 6:30-8PM**
  - All topics up to and including Lecture 15
    - » Focus will be on Lectures 16 – 23 and associated readings, and Projects 3
    - » But expect 20-30% questions from materials from Lectures 1-15
  - A-L 245 Li Ka Shing, M-S 2060 VLSB, T-Z 2040 VLSB
  - Closed book
  - A single hand-written note, both sides

4/19/17

CS162 ©UCB Spring 2017

Lec 23.10

BREAK

4/19/17

CS162 ©UCB Spring 2017

Lec 23.11

### Key Value Storage

---

- Handle huge volumes of data, e.g., PetaBytes!
  - Store (key, value) tuples
- Simple interface
  - `put(key, value);` // insert/write “value” associated with “key”
  - `value = get(key);` // get/read data associated with “key”
- Used sometimes as a simpler but more scalable “database”

4/19/17

CS162 ©UCB Spring 2017

Lec 23.12

## Key Values: Examples

- Amazon:  
– Key: customerID  
– Value: customer profile (e.g., buying history, credit card, ..)
- Facebook, Twitter:  
– Key: UserID  
– Value: user profile (e.g., posting history, photos, friends, ...)
- iCloud/iTunes:  
– Key: Movie/song name  
– Value: Movie, Song



4/19/17

CS162 ©UCB Spring 2017

Lec 23.13

## Key-Value Storage Systems in Real Life

- **Amazon**
  - DynamoDB: internal key value store used for Amazon.com (cart)
  - Simple Storage System (S3)
- **BigTable/HBase/Hypertable**: distributed, scalable data store
- **Cassandra**: “distributed data management system” (developed by Facebook)
- **Memcached**: in-memory key-value store for small chunks of arbitrary data (strings, objects)
- **BitTorrent distributed file location**: peer-to-peer sharing system
- ...


4/19/17

CS162 ©UCB Spring 2017

Lec 23.14

## Key Value Store

- Also called Distributed Hash Tables (DHT)
- Main idea: partition set of key-values across many machines



4/19/17

CS162 ©UCB Spring 2017

Lec 23.15

## Challenges



- **Fault Tolerance**: handle machine failures without losing data and without degradation in performance
- **Scalability**:
  - Need to scale to thousands of machines
  - Need to allow easy addition of new machines
- **Consistency**: maintain data consistency in face of node failures and message losses
- **Heterogeneity** (if deployed as peer-to-peer systems):
  - Latency: 1ms to 1000ms
  - Bandwidth: 32Kb/s to 100Mb/s

4/19/17

CS162 ©UCB Spring 2017

Lec 23.16

## Key Questions

- **put(key, value)**: where to store a new (key, value) tuple?
- **get(key)**: where is the value associated with a given "key" stored?
- And, do the above while providing
  - Fault Tolerance
  - Scalability
  - Consistency

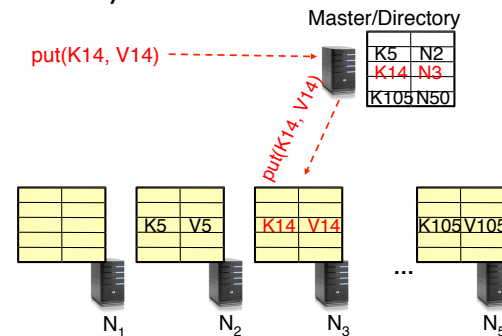
4/19/17

CS162 ©UCB Spring 2017

Lec 23.17

## Directory-Based Architecture

- Have a node maintain the mapping between **keys** and the **machines (nodes)** that store the **values** associated with the **keys**



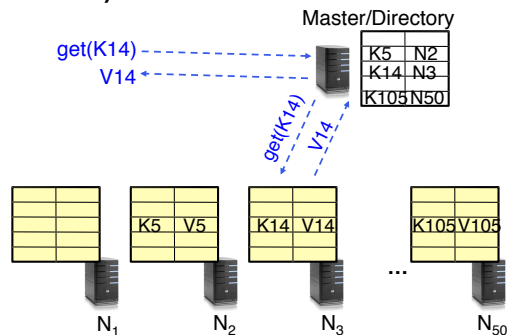
4/19/17

CS162 ©UCB Spring 2017

Lec 23.18

## Directory-Based Architecture

- Have a node maintain the mapping between **keys** and the **machines (nodes)** that store the **values** associated with the **keys**



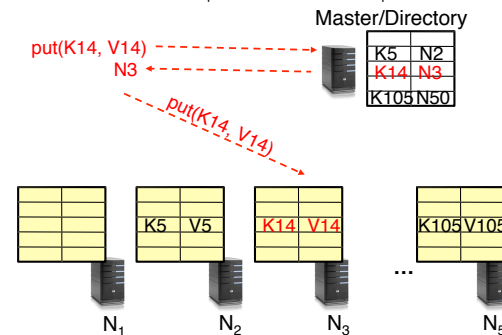
4/19/17

CS162 ©UCB Spring 2017

Lec 23.19

## Directory-Based Architecture

- Having the master relay the requests → **recursive query**
- Another method: **iterative query** (this slide)
  - Return node to requester and let requester contact node



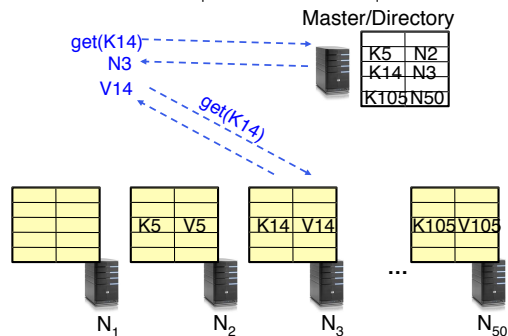
4/19/17

CS162 ©UCB Spring 2017

Lec 23.20

## Directory-Based Architecture

- Having the master relay the requests → **recursive query**
- Another method: **iterative query**
  - Return node to requester and let requester contact node




4/19/17

CS162 ©UCB Spring 2017

Lec 23.21

## Discussion: Iterative vs. Recursive Query



- Recursive Query:
  - Advantages:
    - » Faster, as typically master/directory closer to nodes
    - » Easier to maintain consistency, as master/directory can serialize puts()/gets()
  - Disadvantages: scalability bottleneck, as all “Values” go through master/directory
- Iterative Query
  - Advantages: more scalable
  - Disadvantages: slower, harder to enforce data consistency

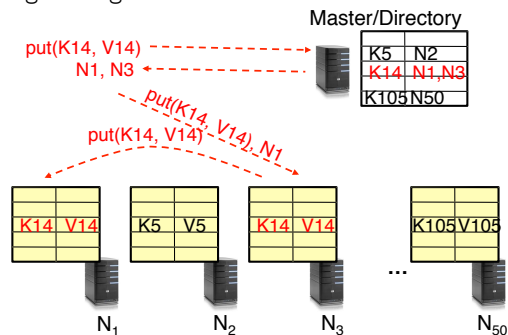
4/19/17

CS162 ©UCB Spring 2017

Lec 23.22

## Fault Tolerance

- Replicate value on several nodes
- Usually, place replicas on different racks in a datacenter to guard against rack failures



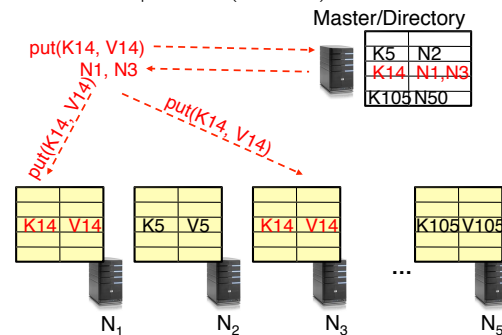
4/19/17

CS162 ©UCB Spring 2017

Lec 23.23

## Fault Tolerance

- Again, we can have
  - **Recursive** replication (previous slide)
  - **Iterative** replication (this slide)



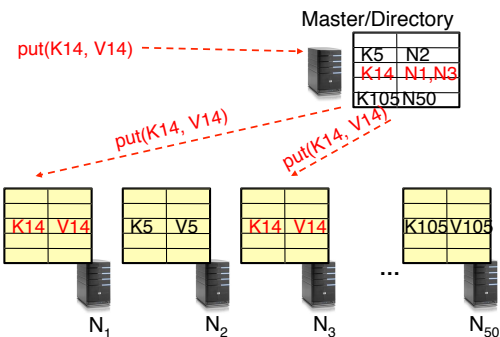
4/19/17

CS162 ©UCB Spring 2017

Lec 23.24

## Fault Tolerance

- Or we can use **recursive** query and **iterative** replication...



4/19/17

CS162 ©UCB Spring 2017

Lec 23.25

## Scalability

- More Storage: use more nodes
- More Requests:
  - Can serve requests from all nodes on which a value is stored in parallel
  - Master can replicate a popular value on more nodes
- Master/directory scalability:
  - Replicate it
  - Partition it, so different keys are served by different masters/directories
    - » How do you partition?

4/19/17

CS162 ©UCB Spring 2017

Lec 23.26

## Scalability: Load Balancing

- Directory keeps track of the storage availability at each node
  - Preferentially insert new values on nodes with more storage available
- What happens when a new node is added?
  - Cannot insert only new values on new node. Why?
  - Move values from the heavy loaded nodes to the new node
- What happens when a node fails?
  - Need to replicate values from fail node to other nodes

4/19/17

CS162 ©UCB Spring 2017

Lec 23.27

## Consistency

- Need to make sure that a value is replicated correctly
- How do you know a value has been replicated on every node?
  - Wait for acknowledgements from every node
- What happens if a node fails during replication?
  - Pick another node and try again
- What happens if a node is slow?
  - Slow down the entire put()? Pick another node?
- In general, with multiple replicas
  - Slow puts and fast gets

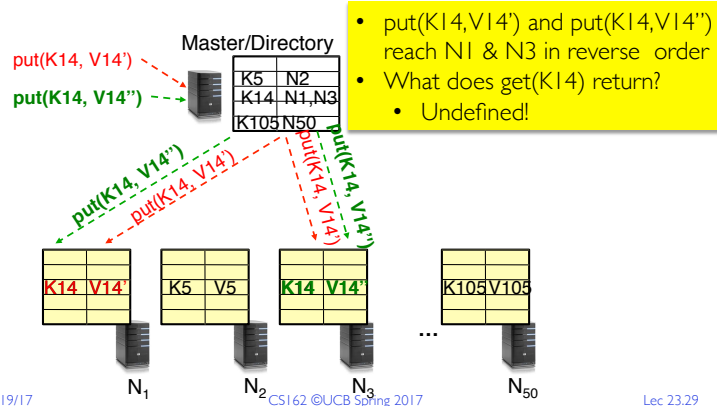
4/19/17

CS162 ©UCB Spring 2017

Lec 23.28

## Consistency (cont'd)

- If concurrent updates (i.e., puts to same key) may need to make sure that updates happen in the same order



## Large Variety of Consistency Models

- Atomic consistency (linearizability): reads/writes (gets/puts) to replicas appear as if there was a single underlying replica (single system image)
  - Think “one updated at a time”
  - Transactions
- Eventual consistency: given enough time all updates will propagate through the system
  - One of the weakest form of consistency; used by many systems in practice
  - Must eventually converge on single value/key (coherence)
- And many others: causal consistency, sequential consistency, strong consistency, ...

4/19/17

CS162 ©UCB Spring 2017

Lec 23.30

## Quorum Consensus

- Improve  $put()$  and  $get()$  operation performance
- Define a replica set of size  $N$ 
  - $put()$  waits for acknowledgements from at least  $W$  replicas
  - $get()$  waits for responses from at least  $R$  replicas
  - $W+R > N$
- Why does it work?
  - There is at least one node that contains the update
- Why might you use  $W+R > N+1$ ?


4/19/17

CS162 ©UCB Spring 2017

Lec 23.31

## Quorum Consensus Example

- $N=3$ ,  $W=2$ ,  $R=2$
- Replica set for  $K14$ :  $\{N_1, N_3, N_4\}$
- Assume  $put()$  on  $N_3$  fails



4/19/17


CS162 ©UCB Spring 2017

Lec 23.32



## Quorum Consensus Example

- Now, issuing `get()` to any two nodes out of three will return the answer



4/19/17

CS162 ©UCB Spring 2017

Lec 23.33

## Scaling Up Directory

- Challenge:
  - Directory contains a number of entries equal to number of (key, value) tuples in the system
  - Can be tens or hundreds of billions of entries in the system!
- Solution: **consistent hashing**
- Associate to each node a unique *id* in a uni-dimensional space  $0..2^m-1$ 
  - Partition this space across  $m$  machines
  - Assume keys are in same uni-dimensional space
  - Each (Key, Value) is stored at the node with the smallest ID larger than Key


4/19/17

CS162 ©UCB Spring 2017

Lec 23.34

## Key to Node Mapping Example

- $m = 8 \rightarrow$  ID space:  $0..63$
- Node 8 maps keys  $[5, 8]$
- Node 15 maps keys  $[9, 15]$
- Node 20 maps keys  $[16, 20]$
- ...
- Node 4 maps keys  $[59, 4]$



4/19/17

CS162 ©UCB Spring 2017

Lec 23.35

## Scaling Up Directory

- With consistent hashing, directory contains only a number of entries equal to number of nodes
  - Much smaller than number of tuples
- Next challenge: every query still needs to contact the directory
- Solution: distributed directory (a.k.a. lookup) service:
  - Given a **key**, find the **node** storing value associated to the key
- Key idea: route request from node to node until reaching the node storing the request's key
- Key advantage: totally distributed
  - No point of failure; no hot spot

4/19/17

CS162 ©UCB Spring 2017

Lec 23.36

## Chord: Distributed Lookup (Directory) Service

- Key design decision
  - Decouple correctness from efficiency
- Properties
  - Each node needs to know about  $O(\log(M))$ , where  $M$  is the total number of nodes
  - Guarantees that a tuple is found in  $O(\log(M))$  steps
- Many other lookup services: CAN, Tapestry, Pastry, Kademlia, ...


4/19/17

CS162 ©UCB Spring 2017

Lec 23.37

## Lookup

- Each node maintains pointer to its successor
- Route packet (Key, Value) to the node responsible for ID using successor pointers
- Eg., node=4 looks up for node responsible for Key=37



4/19/17

CS162 ©UCB Spring 2017

Lec 23.38

## Stabilization Procedure

- Periodic operation performed by each node  $n$  to maintain its successor when new nodes join the system

```

n.stabilize()
  x = succ.pred;
  if (x ∈ (n, succ))
    succ = x; // if x better successor, update
  succ.notify(n); // n tells successor about itself

n.notify(n')
  if (pred = nil or n' ∈ (pred, n))
    pred = n'; // if n' is better predecessor, update
    
```


4/19/17

CS162 ©UCB Spring 2017

Lec 23.39

## Joining Operation

- Node with id=50 joins the ring
- Node 50 needs to know at least one node already in the system
  - Assume known node is 15



4/19/17

CS162 ©UCB Spring 2017

Lec 23.40

### Joining Operation

- n=50 sends join(50) to node 15
- n=44 returns node 58
- n=50 updates its successor to 58

4/19/17 CS162 ©UCB Spring 2017 Lec 23.41

### Joining Operation

- n=50 executes stabilize()
- n's successor (58) returns x = 44

```

n.stabilize()
x = succ.pred;
if (x ∈ (n, succ))
    succ = x;
succ.notify(n);
  
```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.42

### Joining Operation

- n=50 executes stabilize()
  - x = 44
  - succ = 58

```

n.stabilize()
x = succ.pred;
if (x ∈ (n, succ))
    succ = x;
succ.notify(n);
  
```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.43

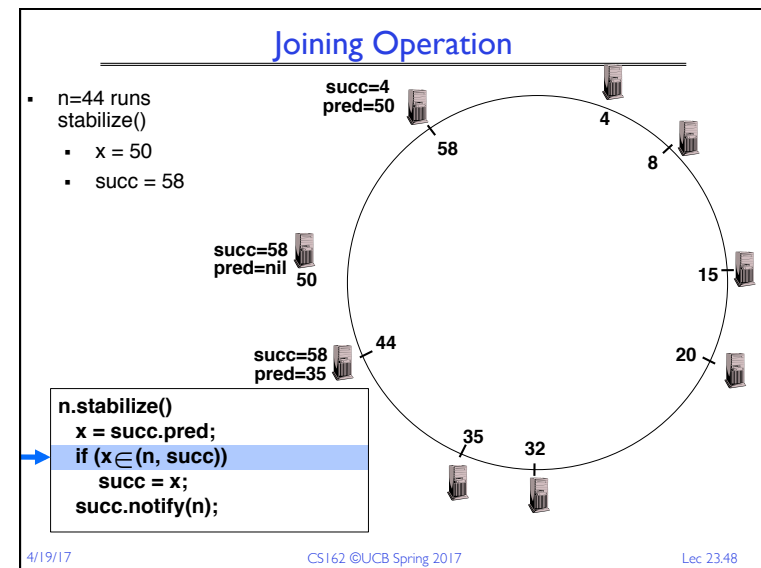
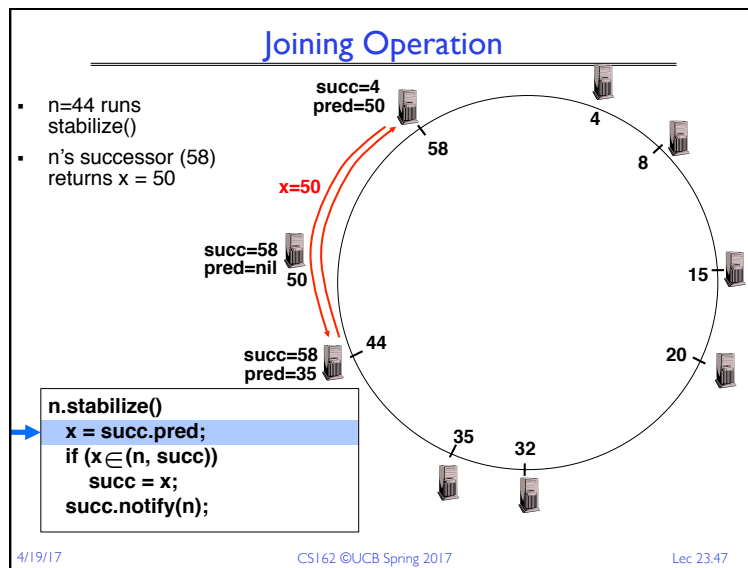
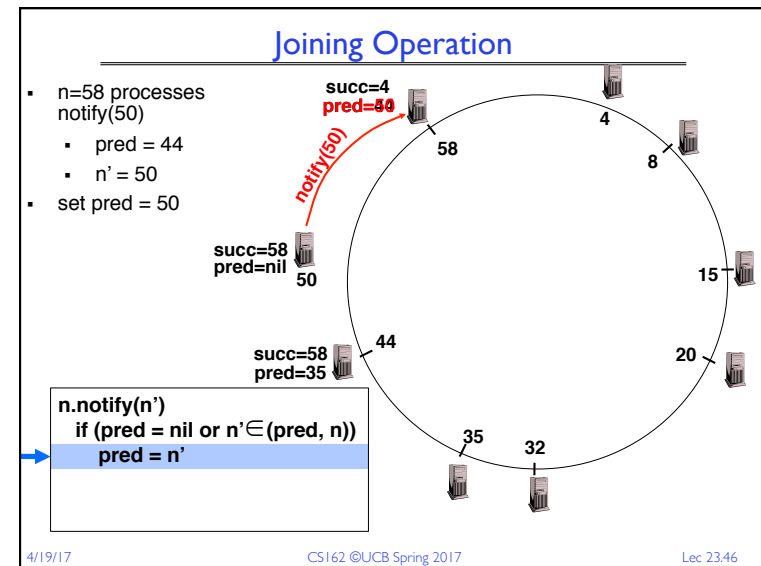
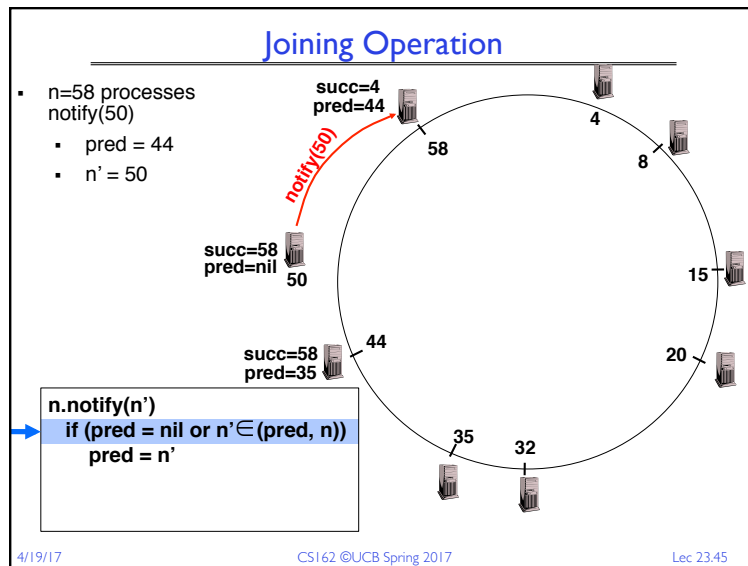
### Joining Operation

- n=50 executes stabilize()
  - x = 44
  - succ = 58
- n=50 sends to its successor (58) notify(50)

```

n.stabilize()
x = succ.pred;
if (x ∈ (n, succ))
    succ = x;
succ.notify(n);
  
```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.44



### Joining Operation

- n=44 runs stabilize()
  - x = 50
  - succ = 58
- n=44 sets succ=50

```

n.stabilize()
x = succ.pred;
if (x ∈ (n, succ))
  succ = x;
  succ.notify(n);

```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.49

### Joining Operation

- n=44 runs stabilize()
- n=44 sends notify(44) to its successor

```

n.stabilize()
x = succ.pred;
if (x ∈ (n, succ))
  succ = x;
  succ.notify(n);

```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.50

### Joining Operation

- n=50 processes notify(44)
  - pred = nil

```

n.notify(n')
if (pred = nil or n' ∈ (pred, n))
  pred = n'

```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.51

### Joining Operation

- n=50 processes notify(44)
  - pred = nil
- n=50 sets pred=44

```


n.notify(n')
if (pred = nil or n' ∈ (pred, n))
  pred = n'

```

4/19/17 CS162 ©UCB Spring 2017 Lec 23.52

## Joining Operation (cont'd)

- This completes the joining operation!



4/19/17


CS162 ©UCB Spring 2017

Lec 23.53

## Achieving Efficiency: *finger tables*

### Finger Table at 80

$i$	$ft[i]$
0	96
1	96
2	96
3	96
4	96
5	112
6	20



$i$ th entry at peer with id  $n$  is first peer with id  $\geq n + 2^i \pmod{2^m}$

4/19/17

CS162 ©UCB Spring 2017

Lec 23.54

## Achieving Fault Tolerance for Lookup Service

- To improve robustness each node maintains the  $k$  ( $> 1$ ) immediate successors instead of only one successor
- In the `pred()` reply message, node A can send its  $k-1$  successors to its predecessor B
- Upon receiving `pred()` message, B can update its successor list by concatenating the successor list received from A with its own list
- If  $k = \log(M)$ , lookup operation works with high probability even if half of nodes fail, where  $M$  is number of nodes in the system


4/19/17

CS162 ©UCB Spring 2017

Lec 23.55

## Storage Fault Tolerance

- Replicate tuples on successor nodes
- Example: replicate (K14, V14) on nodes 20 and 32




4/19/17

CS162 ©UCB Spring 2017

Lec 23.56

## Storage Fault Tolerance

- If node 15 fails, no reconfiguration needed
  - Still have two replicas
  - All lookups will be correctly routed
- Will need to add a new replica on node 35



4/19/17

CS162 ©UCB Spring 2017

Lec 23.57

## Summary (1/2)

- **Remote Procedure Call (RPC):** Call proc on remote machine
  - Provides same interface as procedure
  - Automatic packing and unpacking of arguments (in stub)
- **Key-Value Store:**
  - Two operations
    - » `put(key, value)`
    - » `value = get(key)`
  - Challenges
    - » Fault Tolerance → replication
    - » Scalability → serve `get()`'s in parallel; replicate/cache hot tuples
    - » Consistency → quorum consensus to improve `put()` performance

4/19/17

CS162 ©UCB Spring 2017

Lec 23.58

## Summary (2/2)

- Chord:
  - Highly scalable distributed lookup protocol
  - Each node needs to know about  $O(\log(M))$ , where  $m$  is the total number of nodes
  - Guarantees that a tuple is found in  $O(\log(M))$  steps
  - Highly resilient: works with high probability even if half of nodes fail

4/19/17

CS162 ©UCB Spring 2017

Lec 23.59