

CS162
Operating Systems and
Systems Programming
Lecture 24

Berkeley Data Analytics Stack (BDAS)

April 26th, 2017
Prof. Ion Stoica
<http://cs162.eecs.Berkeley.edu>

Data Deluge

- Billions of users connected through the net
 - WWW, FB, twitter, cell phones, ...
 - 80% of the data on FB was produced last year
 - FB building Exabyte ($2^{60} \approx 10^{18}$) data centers



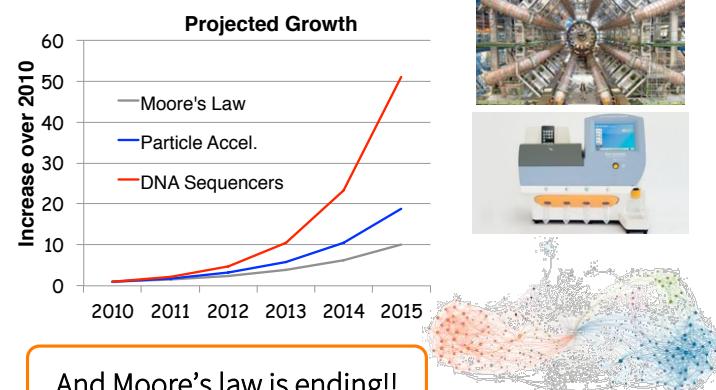
- It's all happening online – could record every:
 - Click, ad impression, billing event, server request, transaction, network msg, fault, fast forward, pause, skip, ...
- User Generated Content (Web & Mobile)
 - Facebook, Instagram, Yelp, TripAdvisor, Twitter, YouTube, ...

4/26/17

CS162 ©UCB Spring 2017

Lec 24.2

Data Grows Faster than Moore's Law



4/26/17

CS162 ©UCB Spring 2017

Lec 24.3

The Big Data Solution: Cloud Computing

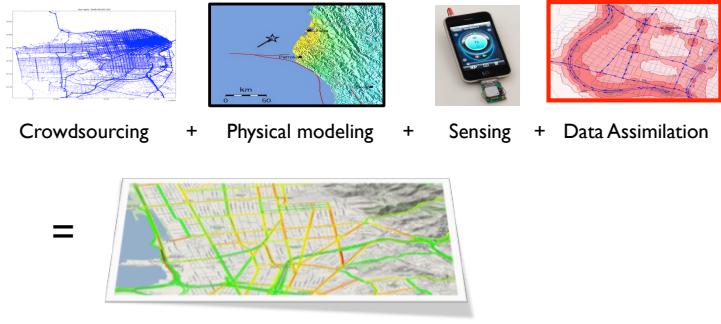
- One machine can not process or even store all the data!
- Solution: distribute data over cluster of cheap machines
 - Lots of hard drives
 - ... and CPUs
 - ... and memory!
- Cloud Computing provides:
 - Illusion of infinite resources
 - Short-term, on-demand resource allocation
 - Can be much less expensive than owning computers
 - Access to latest technologies (SSDs, GPUs, ...)

4/26/17

CS162 ©UCB Spring 2017

Lec 24.4

What Can You do with Big Data?



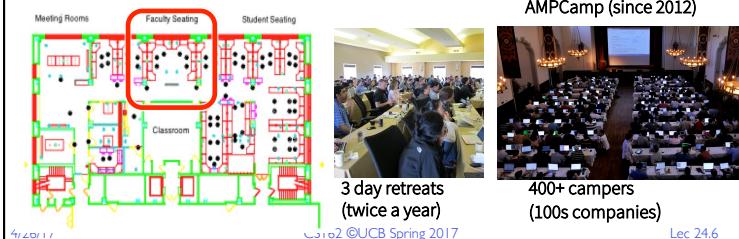
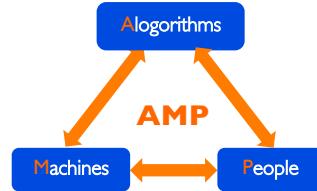
4/26/17

CS162 ©UCB Spring 2017

Lec 24.5

The Berkeley AMPLab

- January 2011 – 2016
 - 8 faculty
 - > 50 students
 - 3 software engineer team
- Organized for collaboration



Lec 24.6

The Berkeley AMPLab

- Governmental and industrial funding:



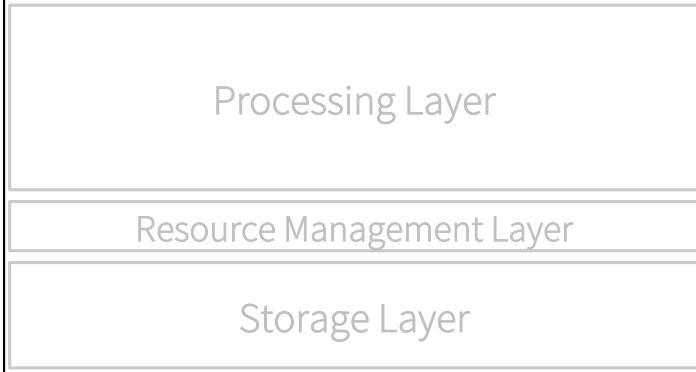
Goal: Next generation of open source data analytics stack for industry & academia:
Berkeley Data Analytics Stack (BDAS)

4/26/17

CS162 ©UCB Spring 2017

Lec 24.7

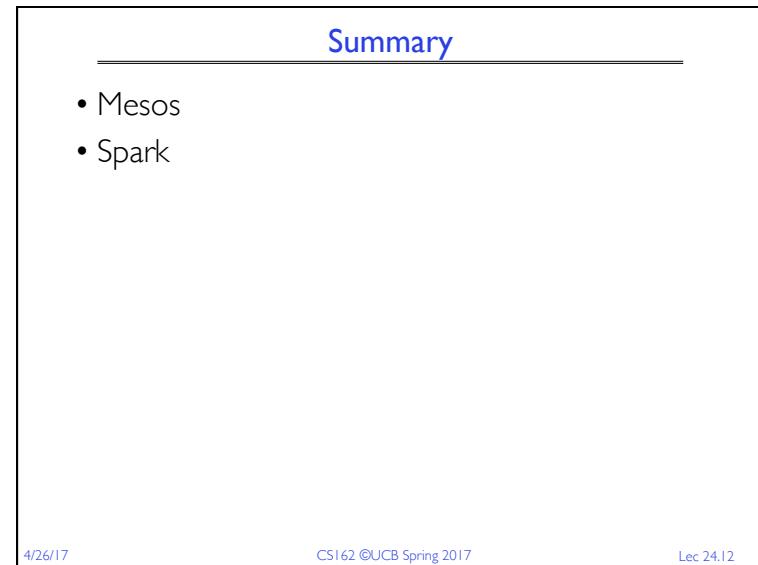
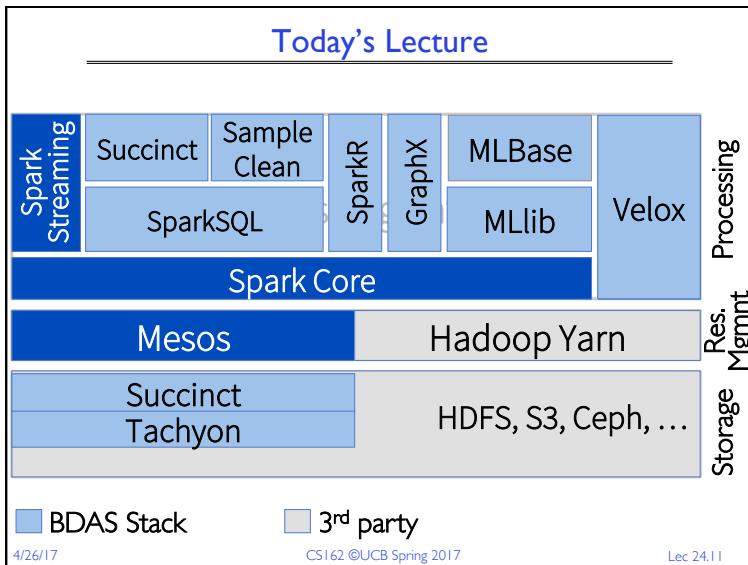
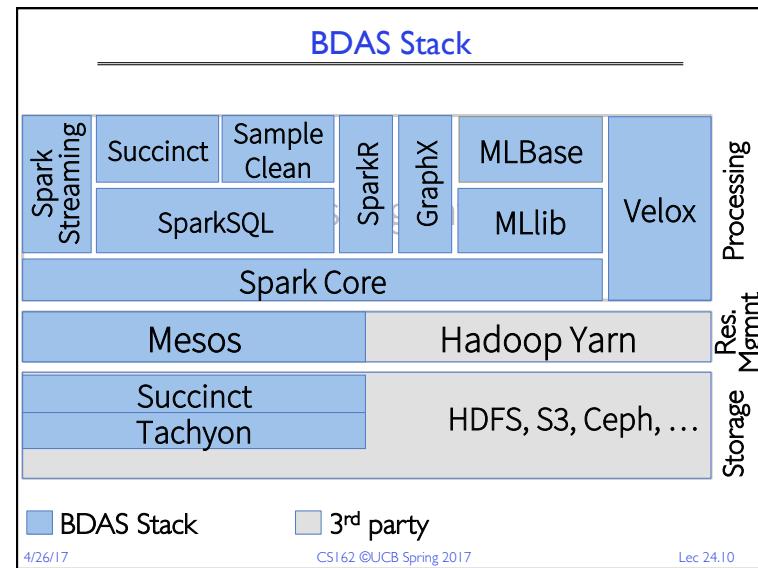
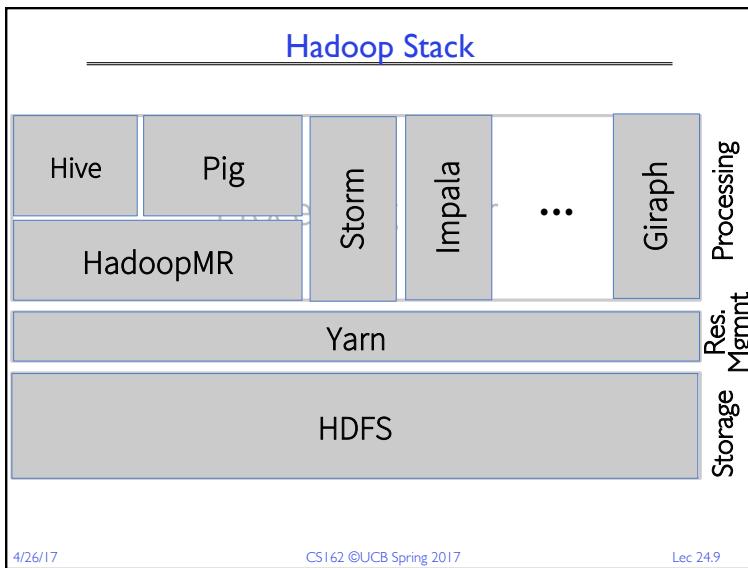
Generic Big Data Stack



4/26/17

CS162 ©UCB Spring 2017

Lec 24.8



Summary

- Mesos
- Spark

4/26/17

CS162 ©UCB Spring 2017

Lec 24.13



A Short History

- Started at UC Berkeley in Spring 2009
 - A class project of [cs294](#) (Cloud Computing: Infrastructure, Services, and Applications)
- Open Source: 2010
- Apache Project: 2011
- Today: one of the most popular cluster resource management systems (OS for datacenters)

4/26/17

CS162 ©UCB Spring 2017

Lec 24.14

Motivation

- Rapid innovation in cloud computing (aka 2008)



- No single framework optimal for all applications
- Each framework runs on its dedicated cluster or cluster partition

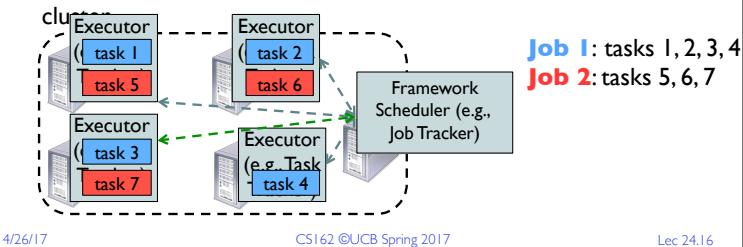
4/26/17

CS162 ©UCB Spring 2017

Lec 24.15

Computation Model: Frameworks

- A **framework** (e.g., Hadoop, MPI) manages one or more **jobs** in a computer cluster
- A **job** consists of one or more **tasks**
- A **task** (e.g., map, reduce) is implemented by one or more processes running on a single machine



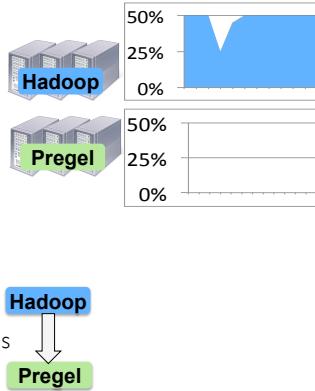
4/26/17

CS162 ©UCB Spring 2017

Lec 24.16

One Framework Per Cluster Challenges

- Inefficient resource usage
 - E.g., Hadoop cannot use available resources from Pregel's cluster
 - No opportunity for stat. multiplexing
- Hard to share data
 - Copy or access remotely, expensive
- Hard to cooperate
 - E.g., Not easy for Pregel to use graphs generated by Hadoop



Need to run multiple frameworks on same cluster

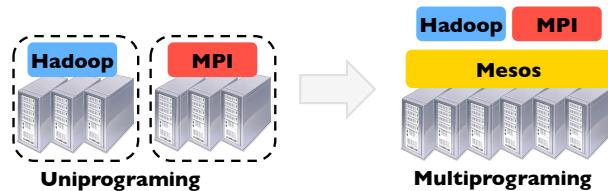
4/26/17

CS162 ©UCB Spring 2017

Lec 24.17

Solution: Apache Mesos

- Common resource sharing layer
 - abstracts ("virtualizes") resources to frameworks
 - enable diverse frameworks to share cluster



4/26/17

CS162 ©UCB Spring 2017

Lec 24.18

Fine Grained Resource Sharing

- Task granularity both in **time & space**
 - Multiplex node/time between tasks belonging to different jobs/frameworks
- Tasks typically short; median $\sim= 10$ sec, minutes
- Why fine grained?
 - Improve data locality
 - Easier to handle node failures

4/26/17

CS162 ©UCB Spring 2017

Lec 24.19

Mesos Goals

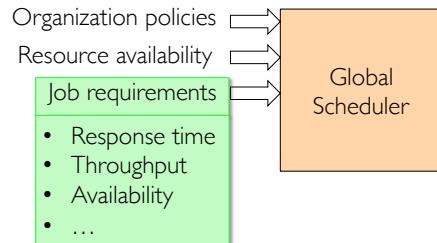
- **High utilization** of resources
- **Support diverse frameworks** (existing & future)
- **Scalability** to 10,000's of nodes
- **Reliability** in face of node failures
- Focus of this talk: **resource management & scheduling**

4/26/17

CS162 ©UCB Spring 2017

Lec 24.20

Approach: Global Scheduler

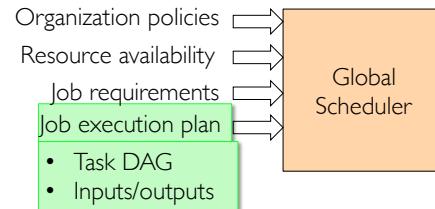


4/26/17

CS162 ©UCB Spring 2017

Lec 24.21

Approach: Global Scheduler

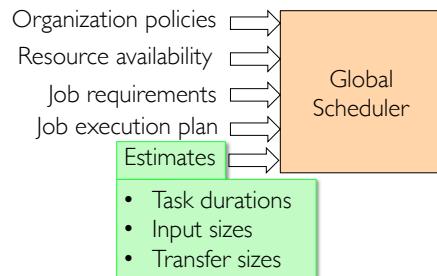


4/26/17

CS162 ©UCB Spring 2017

Lec 24.22

Approach: Global Scheduler

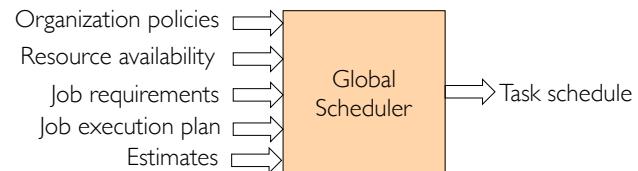


4/26/17

CS162 ©UCB Spring 2017

Lec 24.23

Approach: Global Scheduler



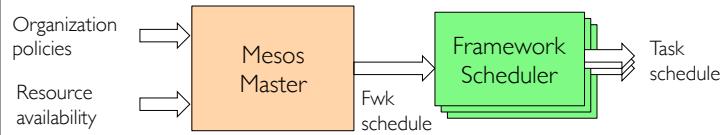
- Advantages: can achieve optimal schedule
- Disadvantages:
 - Complexity → hard to scale and ensure resilience
 - Hard to anticipate future frameworks' requirements
 - Need to refactor existing frameworks

4/26/17

CS162 ©UCB Spring 2017

Lec 24.24

Our Approach: Distributed Scheduler



- Advantages:
 - Simple → easier to scale and make resilient
 - Easy to port existing frameworks, support new ones
- Disadvantages:
 - Distributed scheduling decision → not optimal

4/26/17

CS162 ©UCB Spring 2017

Lec 24.25

Resource Offers

- Unit of allocation: **resource offer**
 - Vector of available resources on a node
 - E.g., node1: <1CPU, 1GB>, node2: <4CPU, 16GB>
- Master sends resource offers to frameworks
- Frameworks select which offers to accept and which tasks to run

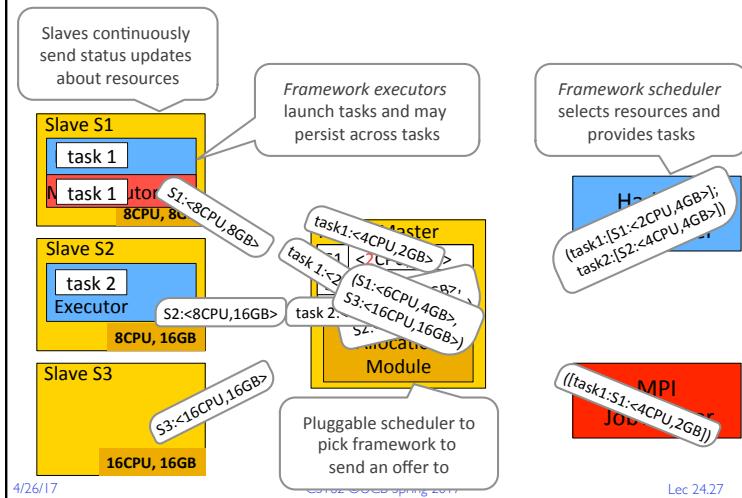
Push task scheduling to frameworks

4/26/17

CS162 ©UCB Spring 2017

Lec 24.26

Mesos Architecture: Example



4/26/17

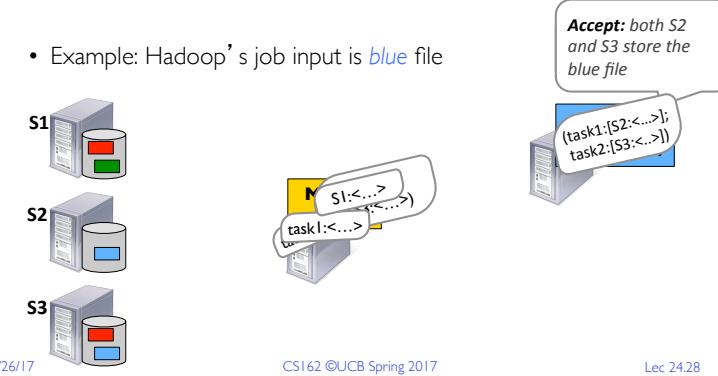
CS162 ©UCB Spring 2017

Lec 24.27

Why does it Work?

- A framework can just wait for an offer that matches its constraints or preferences!
 - **Reject** offers it does not like

- Example: Hadoop's job input is **blue** file



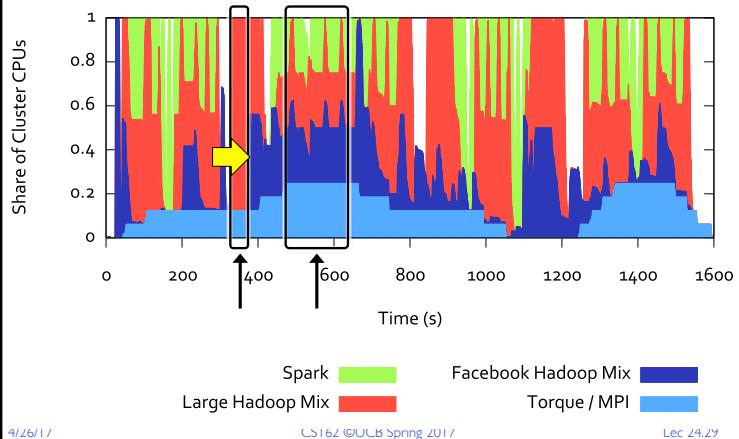
4/26/17

CS162 ©UCB Spring 2017

Lec 24.28

Dynamic Resource Sharing

- 100 node cluster



Apache Mesos Today

- Hundreds of contributors
- Hundreds of deployments in production
 - Eg, Twitter, GE, Apple
 - Managing 10K node datacenters!
- Mesosphere, startup to commercialize Apache Spark

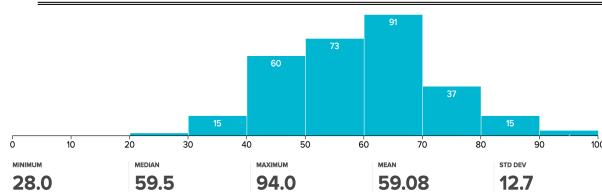


4/26/17

CS162 ©UCB Spring 2017

Lec 24.30

Administrivia



- Midterm #3 grades published
 - Regrade request deadline Wednesday 5/3 at midnight
 - Will open regrading tomorrow, Thursday, 4/27
 - Please only submit regrade requests for grading errors!
- Project 3
 - Code due Monday 5/1
 - Final report due Wednesday 5/5

4/26/17

CS162 ©UCB Spring 2017

Lec 24.31

BREAK

4/26/17

CS162 ©UCB Spring 2017

Lec 24.32

Summary

- Mesos
- Spark

4/26/17

CS162 ©UCB Spring 2017

Lec 24.33

A Short History



- Started at UC Berkeley in 2009
- Open Source: 2010
- Apache Project: 2013
- Today: most popular big data project

4/26/17

CS162 ©UCB Spring 2017

Lec 24.34

What Is Spark?



- Parallel execution engine for big data processing
- Easy to use: 2-5x less code than Hadoop MR
 - High level API's in Python, Java, and Scala
-
- Fast: up to 100x faster than Hadoop MR
 - Can exploit in-memory when available
 - Low overhead scheduling, optimized engine
- General: support multiple computation models

4/26/17

CS162 ©UCB Spring 2017

Lec 24.35

General

- Unifies *batch, interactive* comp.



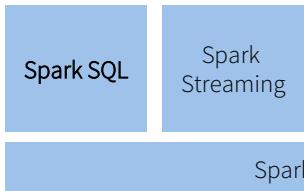
4/26/17

CS162 ©UCB Spring 2017

Lec 24.36

General

- Unifies *batch, interactive, streaming* comp.



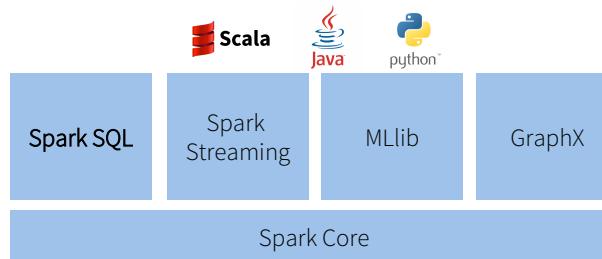
4/26/17

CS162 ©UCB Spring 2017

Lec 24.37

General

- Unifies *batch, interactive, streaming* comp.
- Easy to build sophisticated applications
 - Support iterative, graph-parallel algorithms
 - Powerful APIs in Scala, Python, Java



4/26/17

CS162 ©UCB Spring 2017

Lec 24.38

Easy to Write Code

```
1 public class WordCount {
2     public static void main(String[] args) throws IOException, InterruptedException {
3         Configuration conf = new Configuration();
4         conf.set("mapreduce.job.reduces", "1");
5         Job job = new Job(conf, "word count");
6         job.setMapperClass(WordCountMapper.class);
7         job.setReducerClass(WordCountReducer.class);
8         job.setOutputKeyClass(Text.class);
9         job.setOutputValueClass(IntWritable.class);
10        FileInputFormat.addInputPath(job, new Path(args[0]));
11        FileOutputFormat.setOutputPath(job, new Path(args[1]));
12        System.exit(job.waitForCompletion(true));
13    }
14 }
15
16 public static class WordCountMapper extends Mapper<Text, Text, IntWritable> {
17     private IntWritable one = new IntWritable(1);
18     protected void map(Text key, Text value, Context context)
19     throws IOException, InterruptedException {
20         StringTokenizer itr = new StringTokenizer(value.toString());
21         while (itr.hasMoreTokens()) {
22             context.write(new Text(itr.nextToken()), one);
23         }
24     }
25 }
26
27 public static class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
28     protected void reduce(Text key, IntWritable values, Context context)
29     throws IOException, InterruptedException {
30         int sum = 0;
31         for (IntWritable val : values) {
32             sum += val.get();
33         }
34         context.write(key, new IntWritable(sum));
35     }
36 }
37
38 public static void main(String[] args) throws IOException {
39     Configuration conf = new Configuration();
40     String[] strings = new String[2];
41     strings[0] = args[0];
42     strings[1] = args[1];
43     System.out.println("usage: wordcount <in> [<out>]");
44     Job job = new Job(conf, "word count");
45     job.setMapperClass(WordCountMapper.class);
46     job.setReducerClass(WordCountReducer.class);
47     job.setOutputKeyClass(Text.class);
48     job.setOutputValueClass(IntWritable.class);
49     FileInputFormat.addInputPath(job, new Path(strings[0]));
50     FileOutputFormat.setOutputPath(job, new Path(strings[1]));
51     job.setJarByClass(WordCount.class);
52     job.setNumReduceTasks(1);
53     job.waitForCompletion(true);
54     System.exit(job.waitForCompletion(true));
55 }
```

WordCount in 3 lines of Spark

WordCount in 50+ lines of Java MR

4/26/17

CS162 ©UCB Spring 2017

Lec 24.39

Fast: Time to sort 100TB

2013 Record: 2100 machines
Hadoop



72 minutes



2014 Record:
Spark

207 machines



23 minutes



Also sorted 1PB in 4 hours

4/26/17

Source: Daytona GreySort benchmark, sortbenchmark.org

CS162 ©UCB Spring 2017

Lec 24.40

Large-Scale Usage

Largest cluster: 8000 nodes  Tencent 腾讯

Largest single job: 1 petabyte   Alibaba.com databricks

Top streaming intake: 1 TB/hour  HHMI
janelia farm
research campus

2014 on-disk sort record

4/26/17

CS162 ©UCB Spring 2017

Lec 24.41

RDD: Core Abstraction

Write programs in terms of **distributed datasets**

and **operations** on them

- **Resilient Distributed Datasets (RDDs)**

- Collections of objects distr. across a cluster, stored in RAM or on Disk
- Built through parallel transformations
- Automatically rebuilt on failure

- **Operations**

- Transformations (e.g. map, filter, groupBy)
- Actions (e.g. count, collect, save)

4/26/17

CS162 ©UCB Spring 2017

Lec 24.42

Operations on RDDs

- Transformations $f(\text{RDD}) \Rightarrow \text{RDD}$
 - Lazy (not computed immediately)
 - E.g. "map"
- Actions:
 - Triggers computation
 - E.g. "count", "saveAsTextFile"

4/26/17

CS162 ©UCB Spring 2017

Lec 24.43

Working With RDDs

```
textFile = sc.textFile("SomeFile.txt")
```

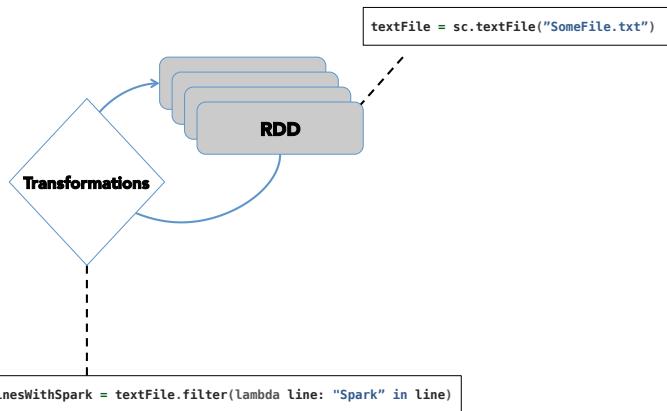
RDD

4/26/17

CS162 ©UCB Spring 2017

Lec 24.44

Working With RDDs

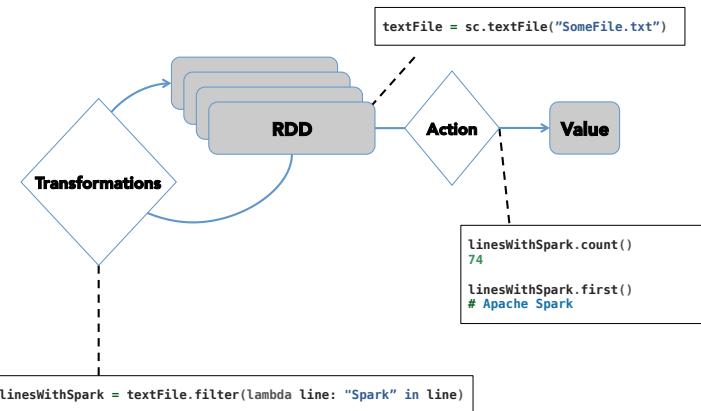


4/26/17

CS162 ©UCB Spring 2017

Lec 24.45

Working With RDDs



4/26/17

CS162 ©UCB Spring 2017

Lec 24.46

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

4/26/17

CS162 ©UCB Spring 2017

Lec 24.47

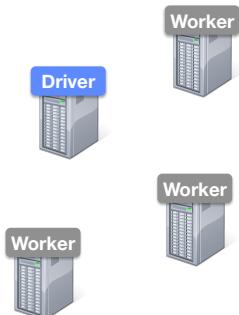
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

4/26/17

CS162 ©UCB Spring 2017

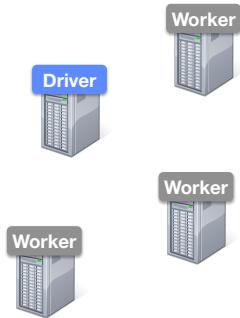
Lec 24.48



Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
```



4/26/17

CS162 ©UCB Spring 2017

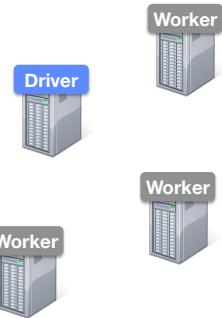
Lec 24.49

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Base RDD

```
lines = spark.textFile("hdfs://...")
```



4/26/17

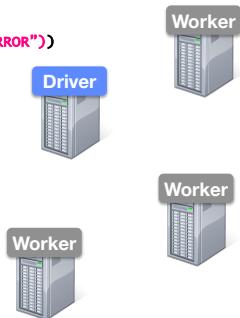
CS162 ©UCB Spring 2017

Lec 24.50

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
```



4/26/17

CS162 ©UCB Spring 2017

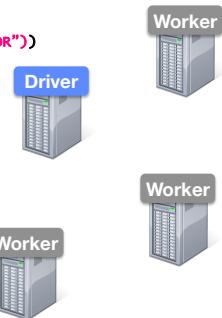
Lec 24.51

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Transformed RDD

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
```



4/26/17

CS162 ©UCB Spring 2017

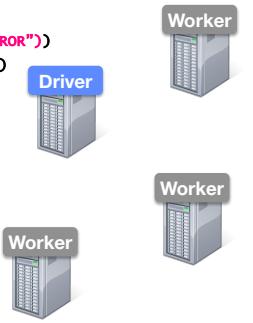
Lec 24.52

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

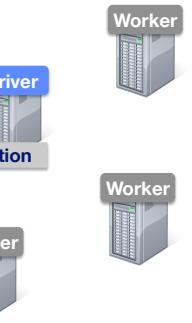
Lec 24.53

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

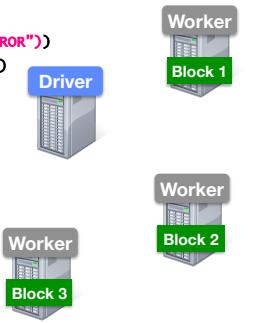
Lec 24.54

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

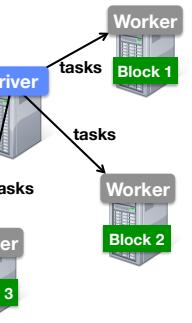
Lec 24.55

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

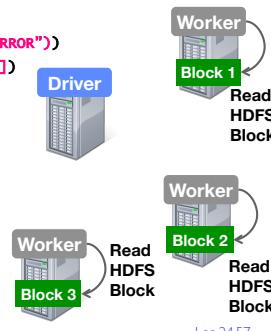
Lec 24.56

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

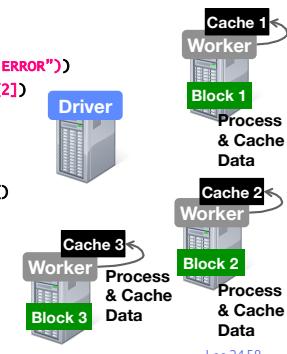
Lec 24.57

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

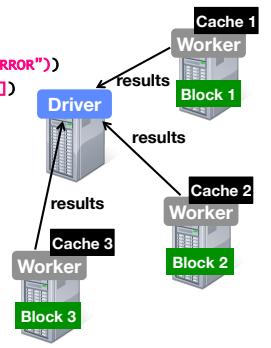
Lec 24.58

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



4/26/17

CS162 ©UCB Spring 2017

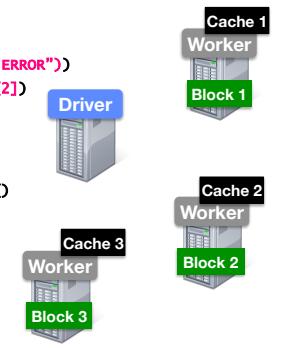
Lec 24.59

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()  
messages.filter(lambda s: "php" in s).count()
```



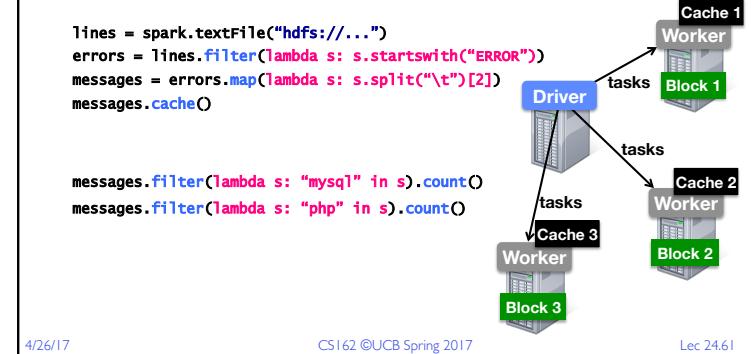
4/26/17

CS162 ©UCB Spring 2017

Lec 24.60

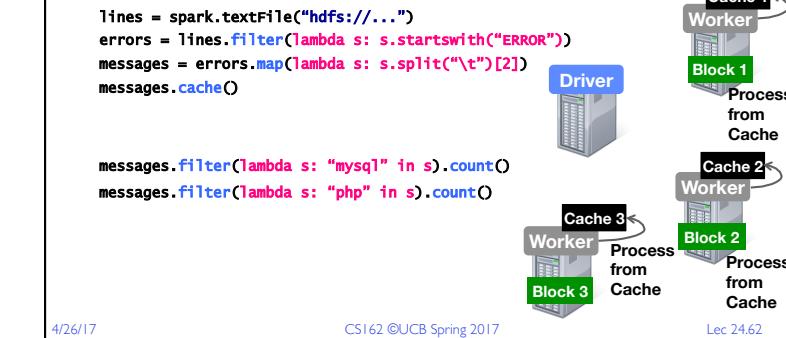
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



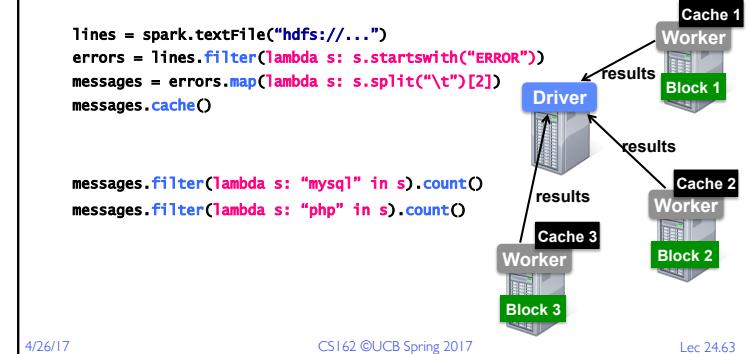
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



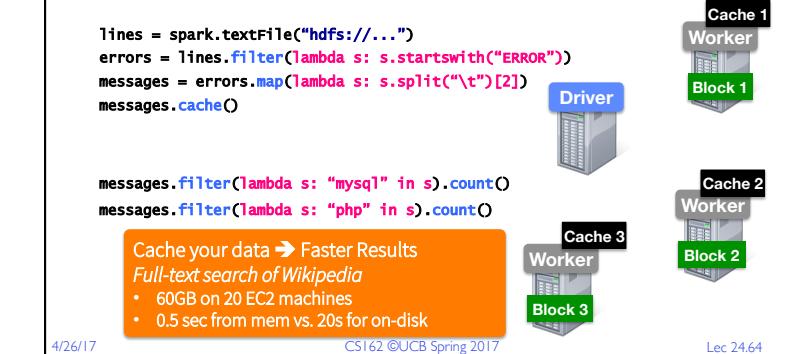
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



Language Support

Python

```
lines = sc.textFile(...)  
lines.filter(lambda s: "ERROR" in s).count()
```

Scala

```
val lines = sc.textFile(...)  
lines.filter(x => x.contains("ERROR")).count()
```

Java

```
JavaRDD<String> lines = sc.textFile(...);  
lines.filter(new Function<String, Boolean>() {  
    Boolean call(String s) {  
        return s.contains("error");  
    }  
}).count();
```

Standalone Programs

- Python, Scala, & Java

Interactive Shells

- Python & Scala

Performance

- Java & Scala are faster due to static typing
- ...but Python is often fine

4/26/17

CS162 ©UCB Spring 2017

Lec 24.65

Expressive API

- map
- reduce

4/26/17

CS162 ©UCB Spring 2017

Lec 24.66

Expressive API

- | | | |
|------------------|-------------|--------------------|
| • map | reduce | sample |
| • filter | count | take |
| • groupBy | fold | first |
| • sort | reduceByKey | partitionBy |
| • union | groupByKey | mapWith |
| • join | cogroup | pipe |
| • leftOuterJoin | cross | save |
| • rightOuterJoin | zip | ... |

4/26/17

CS162 ©UCB Spring 2017

Lec 24.67

Fault Recovery

RDDs track *lineage* information that can be used to efficiently reconstruct lost partitions

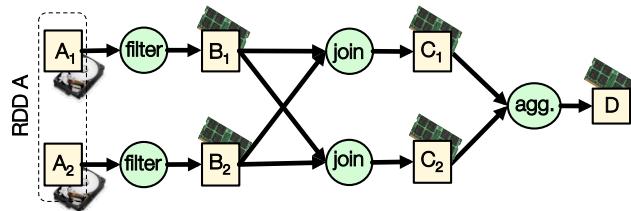
4/26/17

CS162 ©UCB Spring 2017

Lec 24.68

Fault Recovery Example

- Two-partition RDD A={A₁, A₂} stored on disk
 - filter and cache → RDD B
 - join → RDD C
 - aggregate → RDD D



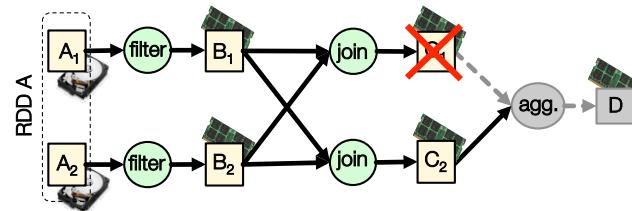
4/26/17

CS162 ©UCB Spring 2017

Lec 24.69

Fault Recovery Example

- C₁ lost due to node failure before reduce finishes



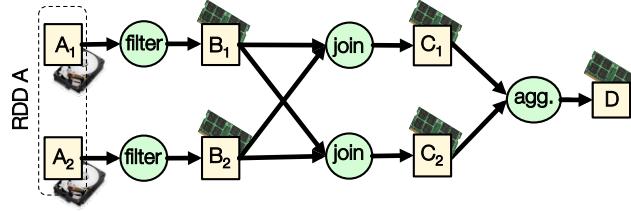
4/26/17

CS162 ©UCB Spring 2017

Lec 24.70

Fault Recovery Example

- C₁ lost due to node failure before reduce finishes
- Reconstruct C₁, eventually, on different node

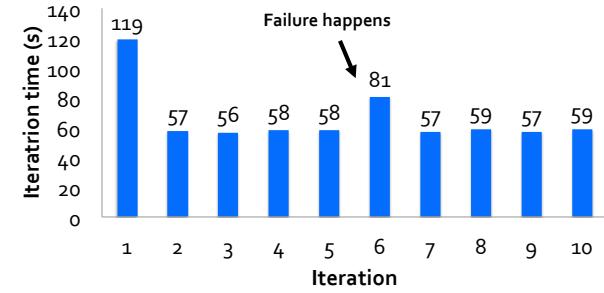


4/26/17

CS162 ©UCB Spring 2017

Lec 24.71

Fault Recovery Example



4/26/17

CS162 ©UCB Spring 2017

Lec 24.72

Spark Streaming: Motivation

- Many important apps must process large data streams at second-scale latencies
 - Site statistics, intrusion detection, online ML
- To build and scale these apps users want:
 - **Integration:** with offline analytical stack
 - **Fault-tolerance:** both for crashes and stragglers

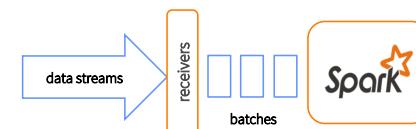
4/26/17

CS162 ©UCB Spring 2017

Lec 24.73

How does it work?

- Data streams are chopped into batches
 - A batch is an RDD holding a few 100s ms worth of data
- Each batch is processed in Spark



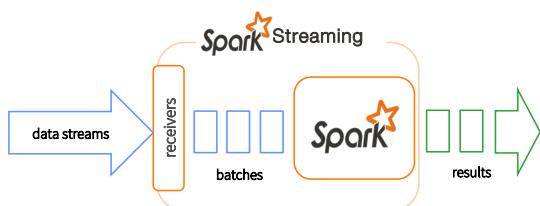
4/26/17

CS162 ©UCB Spring 2017

Lec 24.74

How does it work?

- Data streams are chopped into batches
 - A batch is an RDD holding a few 100s ms worth of data
- Each batch is processed in Spark
- Results pushed out in batches



4/26/17

CS162 ©UCB Spring 2017

Lec 24.75

Streaming Word Count

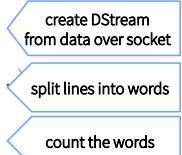
```
val lines = context.socketTextStream("localhost", 9999)
```

```
val words = lines.flatMap(_.split(" "))
```

```
val wordCounts = words.map(x => (x, 1)).reduceByKey(_ + _)
```

```
wordCounts.print()
```

```
ssc.start()
```



start processing the stream

4/26/17

CS162 ©UCB Spring 2017

Lec 24.76

Word Count

```
object NetworkWordCount {
    def main(args: Array[String]) {
        val sparkConf = new SparkConf().setAppName("NetworkWordCount")
        val context = new StreamingContext(sparkConf, Seconds(1))

        val lines = context.socketTextStream("localhost", 9999)
        val words = lines.flatMap(_.split(" "))
        val wordCounts = words.map(x => (x, 1)).reduceByKey(_ + _)

        wordCounts.print()
        ssc.start()
        ssc.awaitTermination()
    }
}
```

4/26/17

CS162 ©UCB Spring 2017

Lec 24.77

Word Count

Spark Streaming

Storm

```

object NetworkWordCount {
    def main(args: Array[String]): Unit = {
        val sparkConf = new SparkConf().setAppName("NetworkWordCount")
        val context = new StreamingContext(sparkConf, Seconds(1))

        val lines = context.socketTextStream("localhost", 9999)
        val words = lines.flatMap(_.split(" "))
        val wordCounts = words.map((_, 1)).reduceByKey(_ + _)

        wordCounts.print()
        sc.start()
        sc.awaitTermination()
    }
}

```

4/26/17

CS162 ©UCB Spring 2017

Lec 24.78

Benefits for Users

- High performance data sharing
 - Data sharing is the bottleneck in many environments
 - RDD's provide in-place sharing through memory
 - Applications can compose models
 - Run a SQL query and then PageRank the results
 - ETL your data and then run graph/ML on it
 - Benefit from investment in shared functionality
 - Eg. re-usable components (shell) and performance optimizations

4/26/17

CS162 ©UCB Spring 2017

Lec 24.7%

Many Recent Developments

- RDDs → Dataframes, and DatSets
 - Distributed collection of data grouped into named columns (i.e. RDD with schema)
 - Think about R and Python Pandas dataframes
 - Project Tungsten
 - Fully managed memory: allowed by Dataframe/DatSets schema
 - Code generation
 - Vectorized processing
 - Structured streams
 - Adding Dataframe API and optimizations to streaming

4/26/17

CS162 ©UCB Spring 2017

Lec 24.80

Apache Spark Today



- > 1,500 contributors
- > 300K committers world wide
- > 100K students trained
- 1,000s deployments in productions
 - Virtually every large enterprise
 - Available in all clouds (e.g., AWS, Google Compute Engine, MS Azure)
 - Distributed by IBM, Cloudera, Hortonworks, Oracle
- Databricks, startup to commercialize Apache Spark



4/26/17

CS162 ©UCB Spring 2017

Lec 24.81

Summary

- Server → Datacenter
- OS → Datacenter OS (e.g., Apache Mesos)
- Applications → Big data / ML applications (e.g., Apache Spark)
- AMPLab
 - Massive success in industry,...
 - and, academia: faculty at MIT, Stanford, Cornell, etc
- New lab starting: RISELab

4/26/17

CS162 ©UCB Spring 2017

Lec 24.82

RISELab

(Real-time Intelligent Secure Execution)



RISELab

From live data to real-time decisions



AMPLab

From batch data to advanced analytics

